

3_2_Models

April 20, 2020

1 SMOTE + FE + Normalization + ML Classification Model

SMOTE → Oversampling technique (called Synthetic Minority Oversampling Technique)

2 1. Import Necessary Libraries

```
[1]: # For Computational and random seed purpose
import numpy as np
np.random.seed(42)
# To read csv file
import pandas as pd
# To Split data into train and cv data
from sklearn.model_selection import train_test_split
# To compute AUROC score
# For AUROC Score (Ref: https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc\_auc\_score.html)
from sklearn.metrics import roc_curve, auc
# Oversampling technique: SMOTE
from imblearn.over_sampling import SMOTE
# Data is unbalance, we need Calibrated Model to give confidence probabilities
↳ result
from sklearn.calibration import CalibratedClassifierCV
# For Hyperparameter and CV Fold
from sklearn.model_selection import GridSearchCV, StratifiedKFold
# For plot AUROC graph
import matplotlib.pyplot as plt
# For heatmap
import seaborn as sns
# To ignore warnings
import warnings
warnings.filterwarnings('ignore')
# To standardize the data
from sklearn.preprocessing import MinMaxScaler
```

```
D:\anaconda3\lib\site-packages\tensorflow\python\framework\dtypes.py:516:
FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,)) /
```

```

'(1,)type'.
_np_qint8 = np.dtype(["qint8", np.int8, 1])
D:\anaconda3\lib\site-packages\tensorflow\python\framework\dtypes.py:517:
FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,)) /
'(1,)type'.
_np_quint8 = np.dtype(["quint8", np.uint8, 1])
D:\anaconda3\lib\site-packages\tensorflow\python\framework\dtypes.py:518:
FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,)) /
'(1,)type'.
_np_qint16 = np.dtype(["qint16", np.int16, 1])
D:\anaconda3\lib\site-packages\tensorflow\python\framework\dtypes.py:519:
FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,)) /
'(1,)type'.
_np_quint16 = np.dtype(["quint16", np.uint16, 1])
D:\anaconda3\lib\site-packages\tensorflow\python\framework\dtypes.py:520:
FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,)) /
'(1,)type'.
_np_qint32 = np.dtype(["qint32", np.int32, 1])
D:\anaconda3\lib\site-packages\tensorflow\python\framework\dtypes.py:525:
FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,)) /
'(1,)type'.
_np_resource = np.dtype(["resource", np.ubyte, 1])
D:\anaconda3\lib\site-packages\tensorboard\compat\tensorflow_stub\dtypes.py:541:
FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,)) /
'(1,)type'.
_np_qint8 = np.dtype(["qint8", np.int8, 1])
D:\anaconda3\lib\site-packages\tensorboard\compat\tensorflow_stub\dtypes.py:542:
FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,)) /
'(1,)type'.
_np_quint8 = np.dtype(["quint8", np.uint8, 1])
D:\anaconda3\lib\site-packages\tensorboard\compat\tensorflow_stub\dtypes.py:543:
FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,)) /
'(1,)type'.
_np_qint16 = np.dtype(["qint16", np.int16, 1])
D:\anaconda3\lib\site-packages\tensorboard\compat\tensorflow_stub\dtypes.py:544:
FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,)) /
'(1,)type'.
_np_quint16 = np.dtype(["quint16", np.uint16, 1])
D:\anaconda3\lib\site-packages\tensorboard\compat\tensorflow_stub\dtypes.py:545:

```

```

FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,)) /
'(1,)type'.
_np_qint32 = np.dtype(["qint32", np.int32, 1])
D:\anaconda3\lib\site-packages\tensorboard\compat\tensorflow_stub\dtypes.py:550:
FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,)) /
'(1,)type'.
np_resource = np.dtype(["resource", np.ubyte, 1])

```

3 2. Read train data

```

[2]: # Locate parent directory
data_dir = "./"

# Read csv file and display top 5 rows
df_train = pd.read_csv(data_dir+'train.csv')
df_train.head(5)

```

```

[2]:   id  target      0      1      2      3      4      5      6      7  ...  \
0    0      1.0 -0.098  2.165  0.681 -0.614  1.309 -0.455 -0.236  0.276  ...
1    1      0.0  1.081 -0.973 -0.383  0.326 -0.428  0.317  1.172  0.352  ...
2    2      1.0 -0.523 -0.089 -0.348  0.148 -0.022  0.404 -0.023 -0.172  ...
3    3      1.0  0.067 -0.021  0.392 -1.637 -0.446 -0.725 -1.035  0.834  ...
4    4      1.0  2.347 -0.831  0.511 -0.021  1.225  1.594  0.585  1.509  ...

      290      291      292      293      294      295      296      297      298      299
0  0.867  1.347  0.504 -0.649  0.672 -2.097  1.051 -0.414  1.038 -1.065
1 -0.165 -1.695 -1.257  1.359 -0.808 -1.624 -0.458 -1.099 -0.936  0.973
2  0.013  0.263 -1.222  0.726  1.444 -1.165 -1.544  0.004  0.800 -1.211
3 -0.404  0.640 -0.595 -0.966  0.900  0.467 -0.562 -0.254 -0.533  0.238
4  0.898  0.134  2.415 -0.996 -1.006  1.378  1.246  1.478  0.428  0.253

[5 rows x 302 columns]

```

```

[3]: df_test = pd.read_csv(data_dir+'test.csv')
df_test.head(5)

```

```

[3]:   id      0      1      2      3      4      5      6      7      8  ...  \
0  250  0.500 -1.033 -1.595  0.309 -0.714  0.502  0.535 -0.129 -0.687  ...
1  251  0.776  0.914 -0.494  1.347 -0.867  0.480  0.578 -0.313  0.203  ...
2  252  1.750  0.509 -0.057  0.835 -0.476  1.428 -0.701 -2.009 -1.378  ...
3  253 -0.556 -1.855 -0.682  0.578  1.592  0.512 -1.419  0.722  0.511  ...
4  254  0.754 -0.245  1.173 -1.623  0.009  0.370  0.781 -1.763 -1.432  ...

      290      291      292      293      294      295      296      297      298      299
0 -0.088 -2.628 -0.845  2.078 -0.277  2.132  0.609 -0.104  0.312  0.979

```

```

1 -0.683 -0.066  0.025  0.606 -0.353 -1.133 -3.138  0.281 -0.625 -0.761
2 -0.094  0.351 -0.607 -0.737 -0.031  0.701  0.976  0.135 -1.327  2.463
3 -0.336 -0.787  0.255 -0.031 -0.836  0.916  2.411  1.053 -1.601 -1.529
4  2.184 -1.090  0.216  1.186 -0.143  0.322 -0.068 -0.156 -1.153  0.825

```

[5 rows x 301 columns]

4 3. Apply Feature Engineering

```

[4]: # We already saw in 2_FE.ipynb file that we created a feat_engg function. We
      ↪ just put it here

def feature_engg(df, if_test = False):
    '''
        Perform Feature Engg in Basic Stats, Trigonometrics, Hyperbolic and
        ↪ Exponential Function

        Parameters:
        df: Pass DataFrame (all features much be in numrric values)
        if_test: If the DataFrame is test data or train data. Ig it is test data,
        ↪ put if_test=True

        Return:
        DataFrame with feature engineering appended
        '''

    if if_test:
        temp = df.drop(['id'], axis=1)
    else:
        temp = df.drop(['id', 'target'], axis=1)

    # Mean and Std FE
    df['mean'] = np.mean(temp, axis=1)
    df['std'] = np.std(temp, axis=1)

    # Trigonometric FE
    sin_temp = np.sin(temp)
    cos_temp = np.cos(temp)
    tan_temp = np.tan(temp)
    df['mean_sin'] = np.mean(sin_temp, axis=1)
    df['mean_cos'] = np.mean(cos_temp, axis=1)
    df['mean_tan'] = np.mean(tan_temp, axis=1)

    # Hyperbolic FE
    sinh_temp = np.sinh(temp)
    cosh_temp = np.cosh(temp)

```

```

tanh_temp = np.tanh(temp)
df['mean_sinh'] = np.mean(sin_temp, axis=1)
df['mean_cosh'] = np.mean(cos_temp, axis=1)
df['mean_tanh'] = np.mean(tan_temp, axis=1)

# Exponents FE
exp_temp = np.exp(temp)
expm1_temp = np.expm1(temp)
exp2_temp = np.exp2(temp)
df['mean_exp'] = np.mean(exp_temp, axis=1)
df['mean_expm1'] = np.mean(expm1_temp, axis=1)
df['mean_exp2'] = np.mean(exp2_temp, axis=1)

# Polynomial FE
# X**2
df['mean_x2'] = np.mean(np.power(temp,2), axis=1)
# X**3
df['mean_x3'] = np.mean(np.power(temp,3), axis=1)
# X**4
df['mean_x4'] = np.mean(np.power(temp,4), axis=1)

return df

```

```

[5]: df_train = feature_engg(df_train)
df_train.head(5)

```

```

[5]:  id  target      0      1      2      3      4      5      6      7  ...  \
0   0      1.0 -0.098  2.165  0.681 -0.614  1.309 -0.455 -0.236  0.276  ...
1   1      0.0  1.081 -0.973 -0.383  0.326 -0.428  0.317  1.172  0.352  ...
2   2      1.0 -0.523 -0.089 -0.348  0.148 -0.022  0.404 -0.023 -0.172  ...
3   3      1.0  0.067 -0.021  0.392 -1.637 -0.446 -0.725 -1.035  0.834  ...
4   4      1.0  2.347 -0.831  0.511 -0.021  1.225  1.594  0.585  1.509  ...

    mean_tan  mean_sinh  mean_cosh  mean_tanh  mean_exp  mean_expm1  mean_exp2  \
0 -0.315591 -0.010536   0.537968 -0.315591  1.760647   0.760647   1.315869
1  0.607457  0.075490   0.611600  0.607457  1.712292   0.712292   1.324817
2  0.104777 -0.005509   0.599358  0.104777  1.749107   0.749107   1.313960
3  0.891722  0.046067   0.645721  0.891722  1.752101   0.752101   1.326229
4  0.274261  0.059548   0.643508  0.274261  1.861741   0.861741   1.377569

    mean_x2  mean_x3  mean_x4
0  1.182425  0.015243  3.584848
1  0.976056  0.047272  2.766570
2  1.023024  0.266454  3.092631
3  0.887980  0.371308  2.553467
4  0.901115  0.613952  2.671541

```

[5 rows x 316 columns]

```
[6]: df_test = feature_engg(df_test, True)
df_test.head(5)
```

```
[6]:
```

	id	0	1	2	3	4	5	6	7	8	...	\
0	250	0.500	-1.033	-1.595	0.309	-0.714	0.502	0.535	-0.129	-0.687	...	
1	251	0.776	0.914	-0.494	1.347	-0.867	0.480	0.578	-0.313	0.203	...	
2	252	1.750	0.509	-0.057	0.835	-0.476	1.428	-0.701	-2.009	-1.378	...	
3	253	-0.556	-1.855	-0.682	0.578	1.592	0.512	-1.419	0.722	0.511	...	
4	254	0.754	-0.245	1.173	-1.623	0.009	0.370	0.781	-1.763	-1.432	...	

	mean_tan	mean_sinh	mean_cosh	mean_tanh	mean_exp	mean_expm1	mean_exp2	\
0	0.565830	0.094378	0.609398	0.565830	1.904397	0.904397	1.404195	
1	-1.641918	-0.018425	0.570495	-1.641918	1.642217	0.642217	1.265487	
2	-0.516155	-0.012641	0.611053	-0.516155	1.517775	0.517775	1.214393	
3	-0.816079	0.002689	0.610619	-0.816079	1.566765	0.566765	1.243412	
4	-1.547172	0.067329	0.611907	-1.547172	1.849024	0.849024	1.374870	

	mean_x2	mean_x3	mean_x4
0	0.985912	0.477020	2.913247
1	1.094274	-0.128315	3.281111
2	0.994294	-0.330590	3.062801
3	0.956136	-0.076546	2.382968
4	0.988710	0.371320	3.079160

[5 rows x 315 columns]

5 4. Split and Oversampling data

```
[7]: # Take separate for features value
X = df_train.drop(['id', 'target'], axis=1)
# Take separate for class value
y = df_train['target'].values
# Take test feature value
ts_X = df_test.drop(['id'], axis=1)
# Split the data into train and cv
tr_X, cv_X, tr_y, cv_y = train_test_split(X, y, test_size=0.1, stratify=y,
    ↪ random_state=42)
# SMOTE (Ref: https://imbalanced-learn.readthedocs.io/en/stable/generated/
    ↪ imblearn.over\_sampling.SMOTE.html)
smote = SMOTE()
# Oversampling using SMOTE technique
tr_X, tr_y = smote.fit_sample(tr_X, tr_y)
```

6 5. Normalization

```
[8]: # Fit and transform on train data
stand_vec = MinMaxScaler()
tr_X = stand_vec.fit_transform(tr_X)
pd.DataFrame(tr_X).head(5)
```

```
[8]:
```

	0	1	2	3	4	5	6	\
0	0.543185	0.385981	0.307849	0.486355	0.593561	0.548421	0.309127	
1	0.536840	0.360374	0.582875	0.362963	0.454180	0.333629	0.769649	
2	0.429595	0.668411	0.542508	0.541326	0.731114	0.284675	0.213147	
3	0.447605	0.656262	0.623445	0.691033	0.192427	0.596488	0.614632	
4	0.161891	0.607477	0.423038	0.435673	0.394366	0.523413	0.904745	

	7	8	9	...	304	305	306	307	\
0	0.730249	0.761705	0.344966	...	0.460236	0.596477	0.145704	0.460236	
1	0.242868	0.711077	0.369924	...	0.326472	0.512572	0.407877	0.326472	
2	0.666423	0.082794	0.230753	...	0.361306	0.051273	0.713982	0.361306	
3	0.600585	0.251618	0.124365	...	0.351714	0.165939	0.542006	0.351714	
4	0.400146	0.468786	0.491751	...	0.335899	0.548545	0.470274	0.335899	

	308	309	310	311	312	313
0	0.758810	0.758810	0.808910	0.755952	0.758091	0.488427
1	0.415875	0.415875	0.454953	0.553698	0.330533	0.472571
2	0.092408	0.092408	0.039308	0.260016	0.274790	0.212496
3	0.169851	0.169851	0.150722	0.347510	0.324467	0.158788
4	0.332553	0.332553	0.382469	0.473672	0.267125	0.345071

[5 rows x 314 columns]

```
[9]: # Transform on cv data on the basis of mean and std generated from train data
cv_X = stand_vec.transform(cv_X)
pd.DataFrame(cv_X).head(5)
```

```
[9]:
```

	0	1	2	3	4	5	6	\
0	0.307818	0.205421	0.471764	0.318129	0.479056	0.361121	0.647591	
1	0.826852	0.712150	0.203466	0.288109	0.579111	0.466300	0.687613	
2	0.199959	0.533832	0.425688	0.530604	0.643314	0.238560	0.845346	
3	0.029267	0.095514	0.616514	0.330994	0.263947	0.682866	0.593263	
4	0.468072	0.616449	0.920082	0.324951	0.685019	0.062433	0.472836	

	7	8	9	...	304	305	306	307	\
0	0.461960	0.555577	0.377538	...	0.364418	0.367169	0.413726	0.364418	
1	0.480432	0.245527	0.666244	...	0.036243	0.691210	0.344602	0.036243	
2	0.520666	1.004949	0.696277	...	-0.026075	0.453845	0.639559	-0.026075	
3	0.445318	0.485535	0.665398	...	0.337094	0.309429	0.414117	0.337094	
4	0.399232	0.474496	0.175127	...	0.252625	0.674429	0.657089	0.252625	

	308	309	310	311	312	313
0	0.313638	0.313638	0.314852	0.522585	0.285660	0.361049
1	0.602803	0.602803	0.665598	0.606822	0.517266	0.476554
2	0.261661	0.261661	0.296326	0.290178	0.354881	0.184931
3	0.392901	0.392901	0.365331	0.565013	0.369783	0.471386
4	0.375732	0.375732	0.443103	0.334018	0.364260	0.308094

[5 rows x 314 columns]

```
[10]: # Transform on test data on the basis of mean and std generated from train data
ts_X = stand_vec.transform(ts_X)
pd.DataFrame(ts_X).head(5)
```

```
[10]:
```

	0	1	2	3	4	5	6	\
0	0.576955	0.354766	0.179817	0.520078	0.338760	0.593650	0.635820	
1	0.633442	0.718692	0.404281	0.722417	0.310774	0.589748	0.643607	
2	0.832788	0.642991	0.493374	0.622612	0.382294	0.757893	0.411988	
3	0.360827	0.201121	0.365953	0.572515	0.760563	0.595424	0.281963	
4	0.628940	0.502056	0.744139	0.143470	0.471008	0.570238	0.680369	

	7	8	9	...	304	305	306	307	\
0	0.455925	0.396079	0.772420	...	0.355177	0.897632	0.535569	0.355177	
1	0.422275	0.565474	0.786168	...	0.310915	0.376359	0.277277	0.310915	
2	0.112107	0.264560	0.534687	...	0.333485	0.403085	0.546558	0.333485	
3	0.611558	0.624096	0.619289	...	0.327472	0.473930	0.543680	0.327472	
4	0.157096	0.254282	0.302665	...	0.312815	0.772634	0.552227	0.312815	

	308	309	310	311	312	313
0	0.777307	0.777307	0.870016	0.418998	0.854945	0.336299
1	0.414707	0.414707	0.422914	0.645653	0.356126	0.446974
2	0.242601	0.242601	0.258222	0.436530	0.189445	0.381293
3	0.310356	0.310356	0.351761	0.356717	0.398786	0.176759
4	0.700724	0.700724	0.775491	0.424850	0.767844	0.386215

[5 rows x 314 columns]

7 6. Apply ML Models (with hyperparameter)

```
[11]: def hyperparameter_model(models, params):
    """
    Hyperparameter tuning with StratifiedKFold follow by GridSearchCV follow by
    ↳ CalibratedClassifier

    Parameters:
    models: Instance of the model
```



```

params: list of parameters with value fr tuning (dict)

Return:
grid_clf: return gridsearch model
'''

# Perform KCrossValidation with stratified target
str_cv = StratifiedKFold(n_splits=10, random_state=42)
# Perform Hyperparamter using GridSearchCV
grid_clf = GridSearchCV(models, params, cv=str_cv, return_train_score=True,
→scoring='roc_auc')
# Fit the train model to evaluate score
grid_clf.fit(tr_X, tr_y)
return grid_clf

# Ref: https://scikit-learn.org/stable/auto\_examples/model\_selection/plot\_roc.html
→html
def plot_roc(try_true, try_pred, cvy_true, cvy_pred, n_classes):
    '''
    Compute ROC curve and ROC area for each class

    Parameters:
    try_true: train true label
    try_pred: train predict probabilities value
    cvy_true: cv true label
    cvy_pred: cv predict probabilities value
    n_classes: number of unique classes

    Return:
    Plot of ROC Curve for train and cv data
    '''
    # For train
    tr_fpr = dict()
    tr_tpr = dict()
    tr_roc_auc = dict()
    for i in range(n_classes):
        tr_fpr[i], tr_tpr[i], _ = roc_curve(try_true, try_pred[:, i])
        tr_roc_auc[i] = auc(tr_fpr[i], tr_tpr[i])

    # For cv
    cv_fpr = dict()
    cv_tpr = dict()
    cv_roc_auc = dict()
    for i in range(n_classes):
        cv_fpr[i], cv_tpr[i], _ = roc_curve(cvy_true, cvy_pred[:, i])
        cv_roc_auc[i] = auc(cv_fpr[i], cv_tpr[i])

    # Line thickness

```

```

lw = 2
# Plot roc for train
plt.plot(tr_fpr[1], tr_tpr[1], color='red',
         lw=lw, label='ROC curve for Train (area = %0.2f)' % tr_roc_auc[1])
# Plot roc for cv
plt.plot(cv_fpr[1], cv_tpr[1], color='green',
         lw=lw, label='ROC curve for CV (area = %0.2f)' % cv_roc_auc[1])
plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic: train vs cv')
plt.legend(loc="lower right")
plt.show()

def plot_feature_importance(model, model_name, top_n = 10):
    """
    Plot the feature importance on the basis of model.

    Parameters:
    model: Instance of model
    model_name: Name of the model
    top_n: Number of feature you want to print top features

    Return:
    df: DataFrame that return feature names with coefficient in descending order
    Plot the feature importance
    """

    # Numpy Column Stack (See Docs: https://docs.scipy.org/doc/numpy-1.10.1/reference/generated/numpy.column\_stack.html)
    column_name = df_train.drop(['id', 'target'], axis=1).columns
    if model_name == 'log_model':
        feat_imp_coef = model.coef_.ravel()
    else:
        feat_imp_coef = model.feature_importances_
    temp = pd.DataFrame(data=np.column_stack((column_name, feat_imp_coef)),
    ↪ columns=['col_name', 'coef'])
    temp = temp.sort_values(by='coef', ascending=False).reset_index()
    df = temp
    temp = temp[:top_n]
    plt.figure(figsize=(20,5))
    sns.barplot(data=temp, y='coef', x='col_name', order=temp['col_name'])
    plt.grid()

```

```

plt.show()
return df

def position_featengg(df):
    '''
    Print the position of feature engg after model fitted

    Parameter:
    df: Pass Dataframe that contain Feature name and their coefficient

    Return:
    Print the rank of the feature engg only!
    '''
    list_feat_engg = [
        'mean', 'std', 'mean_sin', 'mean_cos', 'mean_tan', 'mean_sinh', 'mean_cosh', 'mean_tanh', 'mean_exp',
        'mean_expm1', 'mean_exp2', 'mean_x2', 'mean_x3', 'mean_x4']

    for i in list_feat_engg:
        print('Position rank of', i, ':', df[df['col_name']==i].index[0])

```

7.1 6.1 kNN

```

[14]: # Import KNN
from sklearn.neighbors import KNeighborsClassifier

```

```

[34]: # kNN (See Docs: https://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsClassifier.html)

# List of params
params = {'n_neighbors': np.arange(3, 51, 2).tolist(), 'algorithm': ['kd_tree',
    'brute']}

# Instance of knn model
knn_model = KNeighborsClassifier()

# Call hyperparameter for find the best params as possible
knn_clf = hyperparameter_model(knn_model, params)

```

```

[35]: cv_pvt = pd.pivot_table(pd.DataFrame(knn_clf.cv_results_),
    values='mean_test_score', index='param_n_neighbors', \
        columns='param_algorithm')

tr_pvt = pd.pivot_table(pd.DataFrame(knn_clf.cv_results_),
    values='mean_train_score', index='param_n_neighbors', \
        columns='param_algorithm')


```


```

[36]: plt.title('Train Hyperparameter')
sns.heatmap(tr_pvt, annot=True)
plt.show()

```

```
plt.title('CV Hyperparameter')
sns.heatmap(cv_pvt, annot=True)
plt.show()
```






```
[37]: print(knn_clf.best_params_)
      print('CV Score',knn_clf.score(cv_X,cv_y))

{'algorithm': 'kd_tree', 'n_neighbors': 35}
CV Score 0.7291666666666667

[38]: clf = CalibratedClassifierCV(knn_clf, cv=3)
      clf.fit(tr_X,tr_y)

      tr_pred = clf.predict_proba(tr_X)
      cv_pred = clf.predict_proba(cv_X)

      # Plot ROC cureve of train and cv data
      plot_roc(tr_y, tr_pred, cv_y, cv_pred, 2)
```




8 6.1.1 Kaggle Score

```
[39]: # Create a submission format to make submission in Kaggle
temp_id = df_test['id']
knn_csv = clf.predict_proba(ts_X)[: ,1]
knn_df = pd.DataFrame(np.column_stack((temp_id,knn_csv)),
    ↪columns=['id', 'target'])
knn_df['id'] = knn_df['id'].astype('int32')
knn_df.to_csv(data_dir+'/submission_knn.csv', index=False)
```

```
[40]: image = plt.imread(data_dir+'/submission_knn.png')
plt.figure(figsize=(18,5))
plt.imshow(image)
```

```
[40]: <matplotlib.image.AxesImage at 0x1fd6188b6c8>
```



Observation: Knn perform kind of poorly. So, Knn will not work for this kind of problem

8.1 6.2 Logistic Regression

```
[12]: # Import Logistic Regression
      from sklearn.linear_model import LogisticRegression


[42]: # LogisticRegression (See Docs: https://scikit-learn.org/stable/modules/
      ↪generated/sklearn.linear\_model.LogisticRegression.html)


      # List of hyperparameter that has to be tuned
      params = {'penalty':['l1', 'l2', 'elasticnet'], 'C':[10**i for i in
      ↪range(-4,5)], 'solver':['liblinear', 'sag']}
      # Instance of Logistic Regression
      log_model = LogisticRegression(random_state=42, class_weight='balanced')
      # Call hyperparameter to get the best parameters of this model
      log_clf = hyperparameter_model(log_model, params)

[43]: cv_pvt = pd.pivot_table(pd.DataFrame(log_clf.cv_results_),
      ↪values='mean_test_score', index='param_penalty', \
      columns='param_C')
      tr_pvt = pd.pivot_table(pd.DataFrame(log_clf.cv_results_),
      ↪values='mean_train_score', index='param_penalty', \
      columns='param_C')

[44]: plt.title('Train Hyperparameter')
      sns.heatmap(tr_pvt, annot=True)
      plt.show()

      plt.title('CV Hyperparameter')
      sns.heatmap(cv_pvt, annot=True)
      plt.show()
```






```
[45]: print(log_clf.best_params_)
print('cv Score',log_clf.score(cv_X,cv_y))
```

```
{'C': 1000, 'penalty': 'l1', 'solver': 'liblinear'}
cv Score 0.625
```

```
[46]: clf = CalibratedClassifierCV(log_clf, cv=3)
clf.fit(tr_X,tr_y)

tr_pred = clf.predict_proba(tr_X)
cv_pred = clf.predict_proba(cv_X)


# Plot ROC curve of train and cv data
plot_roc(tr_y, tr_pred, cv_y, cv_pred, 2)
```



```
[47]: # Instance the model passing the best params we got
log_model = LogisticRegression(**log_clf.best_params_, random_state=42,
    ↪class_weight='balanced')
log_model.fit(tr_X, tr_y)
```

```
[47]: LogisticRegression(C=1000, class_weight='balanced', dual=False,
    fit_intercept=True, intercept_scaling=1, l1_ratio=None,
    max_iter=100, multi_class='auto', n_jobs=None, penalty='l1',
    random_state=42, solver='liblinear', tol=0.0001, verbose=0,
    warm_start=False)
```

```
[48]: # Plot the feature importance on the basis of logistic Regression
df = plot_feature_importance(log_model, 'log_model', 20)
```



```
[49]: print('After applying Logistic regression\n')
      position_featengg(df)
```

After applying Logistic regression


```
Position rank of mean : 284
Position rank of std : 120
Position rank of mean_sin : 250
Position rank of mean_cos : 121
Position rank of mean_tan : 122
Position rank of mean_sinh : 259
Position rank of mean_cosh : 123
Position rank of mean_tanh : 124
Position rank of mean_exp : 125
Position rank of mean_expm1 : 126
Position rank of mean_exp2 : 127
Position rank of mean_x2 : 128
Position rank of mean_x3 : 245
Position rank of mean_x4 : 156
```

8.2 6.2.1 Kaggle Score

```
[50]: # Create a submssion format to make submission in Kaggle
      temp_id = df_test['id']
      log_csv = clf.predict_proba(ts_X)[: ,1]
      log_df = pd.DataFrame(np.column_stack((temp_id,log_csv)),
      ↪columns=['id', 'target'])
      log_df['id'] = log_df['id'].astype('int32')
      log_df.to_csv(data_dir+'/submission_log.csv', index=False)
```

```
[51]: image = plt.imread(data_dir+'/submission_log.png')
      plt.figure(figsize=(18,5))
      plt.imshow(image)
```

```
[51]: <matplotlib.image.AxesImage at 0x1fd554e7a08>
```



8.3 6.3 SVC

```
[13]: # Import SVC
      from sklearn.svm import SVC
```



```
[53]: # SVC (See Docs: https://scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html)

      # List of hyperparameter that has to be tuned
      params = {'C': [10**i for i in range(-4,5)], 'kernel':
        ↳ ['linear', 'poly', 'sigmoid', 'rbf']}
      # Instance of SVC
      svc_model = SVC(class_weight='balanced', random_state=42, probability=True)
      # Call hyperparameter to find the best parameters
      svc_clf = hyperparameter_model(svc_model, params)
```

```
[54]: cv_pvt = pd.pivot_table(pd.DataFrame(svc_clf.cv_results_),
        ↳ values='mean_test_score', index='param_C', \
          columns='param_kernel')
      tr_pvt = pd.pivot_table(pd.DataFrame(svc_clf.cv_results_),
        ↳ values='mean_train_score', index='param_C', \
          columns='param_kernel')

      plt.title('Train Hyperparameter')
      sns.heatmap(tr_pvt, annot=True)
      plt.show()

      plt.title('CV Hyperparameter')
      sns.heatmap(cv_pvt, annot=True)
      plt.show()
```




```
[56]: print(svc_clf.best_params_)
print('cv Score',svc_clf.score(cv_X,cv_y))
```

```
{'C': 10, 'kernel': 'rbf'}
cv Score 0.6597222222222222
```

```
[57]: clf = CalibratedClassifierCV(svc_clf, cv=3)
clf.fit(tr_X,tr_y)

tr_pred = clf.predict_proba(tr_X)
cv_pred = clf.predict_proba(cv_X)

# Plot ROC curve of this model
plot_roc(tr_y, tr_pred, cv_y, cv_pred, 2)
```




8.4 6.3.1 Kaggle Score

```
[58]: # Create a submission format to make submission in Kaggle
temp_id = df_test['id']
svc_csv = clf.predict_proba(ts_X)[: ,1]
svc_df = pd.DataFrame(np.column_stack((temp_id,svc_csv)),u
    ↪ columns=['id','target'])
svc_df['id'] = svc_df['id'].astype('int32')
```

```
svc_df.to_csv(data_dir+'/submission_svc.csv', index=False)
```

```
[76]: image = plt.imread(data_dir+'/submission_svc.png')
plt.figure(figsize=(18,5))
plt.imshow(image)
```

```
[76]: <matplotlib.image.AxesImage at 0x1fd55683bc8>
```



8.5 6.4 Random Forest

```
[14]: # Import Random Forest
from sklearn.ensemble import RandomForestClassifier
```



```
[60]: # RandomForest (See Docs: https://scikit-learn.org/stable/modules/generated/
      ↪sklearn.ensemble.RandomForestClassifier.html)

# List of hyperparameter that has to be tuned
params = {'n_estimators': [10, 20, 30, 40, 50, 100, 200, 300, 400], 'max_depth': [2, 3, 5, 7]}
# Instance of randomforest
rf_model = RandomForestClassifier(random_state=42)
# Perform GridSearchCV to find best parameters
rf_clf = hyperparameter_model(rf_model, params)
```

```
[61]: # Ref: https://stackoverflow.com/questions/48791709/
      ↪how-to-plot-a-heat-map-on-pivot-table-after-grid-search

# Plotting of hyperparameter of train and cv score
pvt_tr = pd.pivot_table(pd.DataFrame(rf_clf.cv_results_),
      ↪values='mean_train_score', index='param_n_estimators',
      ↪columns='param_max_depth')
pvt_cv = pd.pivot_table(pd.DataFrame(rf_clf.cv_results_),
      ↪values='mean_test_score', index='param_n_estimators',
      ↪columns='param_max_depth')
plt.figure(1)
plt.title('Hyperparameter for Train')
sns.heatmap(pvt_tr, annot=True)
plt.figure(2)
plt.title('Hyperparameter for CV')
sns.heatmap(pvt_cv, annot=True)
```

```
plt.show()
```




```
[62]: print(rf_clf.best_params_)
```

```
{'max_depth': 5, 'n_estimators': 300}
```

```
[63]: rf_clf = RandomForestClassifier(**rf_clf.best_params_, random_state=42)
      rf_clf.fit(tr_X, tr_y)
```

```
# Calibrate the model
```

```
clf = CalibratedClassifierCV(rf_clf, cv=3)
clf.fit(tr_X, tr_y)
```

```
[63]: CalibratedClassifierCV(base_estimator=RandomForestClassifier(bootstrap=True,
                                                                    ccp_alpha=0.0,
                                                                    class_weight=None,
                                                                    criterion='gini',
                                                                    max_depth=5,

                                                                    max_features='auto',
                                                                    max_leaf_nodes=None,

                                                                    max_samples=None,


                                                                    min_impurity_decrease=0.0,
                                                                    min_impurity_split=None,

                                                                    min_samples_leaf=1,


                                                                    min_samples_split=2,
                                                                    min_weight_fraction_leaf=0.0,

                                                                    n_estimators=300,
                                                                    n_jobs=None,
                                                                    oob_score=False,
                                                                    random_state=42,
                                                                    verbose=0,
                                                                    warm_start=False),
                                                                    cv=3, method='sigmoid')
```

```
[64]: # Plot ROC Curve of train and cv
      plot_roc(tr_y, clf.predict_proba(tr_X), cv_y, clf.predict_proba(cv_X), 2)
```



```
[65]: # Plot the feature importance based on this model
df = plot_feature_importance(rf_clf, 'rf',30)
```



```
[66]: print('After applying Random Forest\n')
position_featengg(df)
```

After applying Random Forest

Position rank of mean : 19

Position rank of std : 81

Position rank of mean_sin : 40

Position rank of mean_cos : 56

```

Position rank of mean_tan : 257
Position rank of mean_sinh : 52
Position rank of mean_cosh : 57
Position rank of mean_tanh : 181
Position rank of mean_exp : 208
Position rank of mean_expm1 : 217
Position rank of mean_exp2 : 73
Position rank of mean_x2 : 172
Position rank of mean_x3 : 108
Position rank of mean_x4 : 313

```

8.6 6.4.1 Kaggle Score

```

[67]: # Create a submission file format to submit in kaggle
temp_id = df_test['id']
rf_csv = clf.predict_proba(ts_X)[: ,1]
rf_df = pd.DataFrame(np.column_stack((temp_id,rf_csv)), columns=['id','target'])
rf_df['id'] = rf_df['id'].astype('int32')
rf_df.to_csv(data_dir+'/submission_rf.csv', index=False)

```

```


[82]: image = plt.imread(data_dir+'/submission_rf.png')
plt.figure(figsize=(18,5))
plt.imshow(image)

```

```

[82]: <matplotlib.image.AxesImage at 0x1fd56b90a88>

```



8.7 6.5 Xgboost

```

[15]: # Import Xgboost
from xgboost import XGBClassifier

```

```

[78]: # Xgboost (See Docs: https://xgboost.readthedocs.io/en/latest/python/python\_api.html)


# List of hyperparameter that has to be tuned
params = {'max_depth':[2,3,5,7], 'n_estimators':[10,20,50,100,200,300,400]}
# Instance of XGBoost Model
xgb_model = XGBClassifier(scale_pos_weight=0.5)
# Call hyperparameter to find the best parameters


```

```
xgb_clf = hyperparameter_model(xgb_model, params)
```

```
[79]: # Ref: https://stackoverflow.com/questions/48791709/how-to-plot-a-heat-map-on-pivot-table-after-grid-search

# Plotting of hyperparameter of train and cv score
pvt_tr = pd.pivot_table(pd.DataFrame(xgb_clf.cv_results_),
    ↪ values='mean_train_score', index='param_n_estimators',
    ↪ columns='param_max_depth')
pvt_cv = pd.pivot_table(pd.DataFrame(xgb_clf.cv_results_),
    ↪ values='mean_test_score', index='param_n_estimators',
    ↪ columns='param_max_depth')
plt.figure(1)
plt.title('Hyperparameter for Train')
sns.heatmap(pvt_tr, annot=True)
plt.figure(2)
plt.title('Hyperparameter for CV')
sns.heatmap(pvt_cv, annot=True)
plt.show()
```






```
[80]: print(xgb_clf.best_params_)
      print('cv Score', xgb_clf.score(cv_X, cv_y))
```

```
{'max_depth': 3, 'n_estimators': 200}
cv Score 0.8125
```


```
[84]: # Instance of randomforest with best parameters
xgb_clf = XGBClassifier(**xgb_clf.best_params_, random_state=42,
    ↪ scale_pos_weight=0.5)
# Fit the model
xgb_clf.fit(tr_X, tr_y)
# Calibrate the model
clf = CalibratedClassifierCV(xgb_clf, cv=3)
clf.fit(tr_X, tr_y)

tr_pred = clf.predict_proba(tr_X)
cv_pred = clf.predict_proba(cv_X)

# Plot ROC curve of train and cv
plot_roc(tr_y, tr_pred, cv_y, cv_pred, 2)
```



```
[85]: # Instance of XGBoost model with best parameters
df = plot_feature_importance(xgb_clf, 'xgb',10)
```



```
[86]: print('After applying Gradient Boosting Random Forest\n')
position_featengg(df)
```

After applying Gradient Boosting Random Forest

Position rank of mean : 67
 Position rank of std : 160
 Position rank of mean_sin : 161
 Position rank of mean_cos : 162

```

Position rank of mean_tan : 29
Position rank of mean_sinh : 163
Position rank of mean_cosh : 164
Position rank of mean_tanh : 165
Position rank of mean_exp : 166
Position rank of mean_expm1 : 167
Position rank of mean_exp2 : 8
Position rank of mean_x2 : 168
Position rank of mean_x3 : 84
Position rank of mean_x4 : 313

```

8.8 6.5.1 Kaggle Score

```

[87]: # Create submission file format to submit in Kaggle
temp_id = df_test['id']
xgb_csv = clf.predict_proba(ts_X)[: ,1]
xgb_df = pd.DataFrame(np.column_stack((temp_id,xgb_csv)),
    ↪columns=['id', 'target'])
xgb_df['id'] = xgb_df['id'].astype('int32')
xgb_df.to_csv(data_dir+'submission_xgb.csv', index=False)

```

```


[88]: image = plt.imread(data_dir+'submission_xgb.png')
plt.figure(figsize=(18,5))
plt.imshow(image)

```

```

[88]: <matplotlib.image.AxesImage at 0x1fd571ee148>

```



8.9 6.6 Stacking Model

```

[16]: # Import Stacking Classifier
from mlxtend.classifier import StackingClassifier

```

```

[17]: # StackClassifier (See Docs: http://rasbt.github.io/mlxtend/user_guide/
    ↪classifier/StackingClassifier/#methods)

# Classifier 1: Logistic Regression with best params
clf1 = LogisticRegression(C = 1000, penalty = 'l1', solver = 'liblinear',
    ↪class_weight='balanced', random_state=42)
clf1.fit(tr_X,tr_y)

```

```

clf1 = CalibratedClassifierCV(clf1, cv=3)

# Classifier 2: SVC with best params
clf2 = SVC(C=10, kernel='rbf', random_state=42, class_weight='balanced',
    ↪probability=True)
clf2.fit(tr_X, tr_y)
clf2 = CalibratedClassifierCV(clf2, cv=3)

# Classifier 3: XGBoost with best params
clf3 = XGBClassifier(max_depth=3, n_estimators=200, scale_pos_weight=0.5)
clf3.fit(tr_X, tr_y)
clf3 = CalibratedClassifierCV(clf3, cv=3)

# Classifier 4: RF with best params
clf4 = RandomForestClassifier(max_depth=5, n_estimators=300)
clf4.fit(tr_X, tr_y)
clf4 = CalibratedClassifierCV(clf4, cv=3)

# Stack Classifier
sclf = StackingClassifier(classifiers=[clf1, clf2, clf3, clf4],
    ↪meta_classifier=clf1, use_proba=True)

# Fit the model
sclf.fit(tr_X, tr_y)

# Predict in probabilities
tr_pred = sclf.predict_proba(tr_X)
cv_pred = sclf.predict_proba(cv_X)

```

```

[91]: # Score after stacking classifier
sclf.score(cv_X, cv_y)


```

[91]: 0.68

```

[92]: # Plot ROC Curve for train and cv
plot_roc(tr_y, tr_pred, cv_y, cv_pred, 2)

```





8.10 6.6.1 Kaggle Score

```
[93]: # Create a submission file format to submit in Kaggle
temp_id = df_test['id']
sclf_csv = sclf.predict_proba(ts_X)[: ,1]
sclf_df = pd.DataFrame(np.column_stack((temp_id,sclf_csv)),
    ↪columns=['id', 'target'])
sclf_df['id'] = sclf_df['id'].astype('int32')
sclf_df.to_csv(data_dir+'/submission_sclf.csv', index=False)
```

```
[94]: image = plt.imread(data_dir+'/submission_sclf.png')
plt.figure(figsize=(18,5))
plt.imshow(image)
```

```
[94]: <matplotlib.image.AxesImage at 0x1fd57201988>
```




9 6.7 Voting Classifier (Without Stack Classifier + no weights)

```
[18]: # Import Voting Classifier
from mlxtend.classifier import EnsembleVoteClassifier
```

```
[19]: # Voting Classifier (See Docs: http://rasbt.github.io/mlxtend/user\_guide/
      ↪ classifier/EnsembleVoteClassifier/)
eclf = EnsembleVoteClassifier(clfs=[clf1, clf2, clf3, clf4])
# Fit the train data
eclf.fit(tr_X, tr_y)
```

```
[19]: EnsembleVoteClassifier(clfs=[CalibratedClassifierCV(base_estimator=LogisticRegression(C=1000,
      class_weight='balanced',
      dual=False,
      fit_intercept=True,
      intercept_scaling=1,
      l1_ratio=None,
      max_iter=100,
      multi_class='auto',
      n_jobs=None,
      penalty='l1',
      random_state=42,
      solver='liblinear',
      tol=0.0001,
      verbose=0,
      warm_start=False),
      cv=3, method='sigmoid'),
      CalibratedClassifierCV(DecisionTreeClassifier(max_depth=5,
      max_features='auto',
      max_leaf_nodes=None,
      max_samples=None,
      min_impurity_decrease=0.0,
      min_impurity_split=None,
      min_samples_leaf=1,
      min_samples_split=2,
      min_weight_fraction_leaf=0.0,
      n_estimators=300,
      n_jobs=None,
      oob_score=False,
      random_state=None,
      verbose=0,
      warm_start=False),
      cv=3, method='sigmoid')],
      refit=True, verbose=0, voting='hard', weights=None)
```

```
[20]: # Predict in probabilities
tr_pred = eclf.predict_proba(tr_X)
cv_pred = eclf.predict_proba(cv_X)
# Plot ROC Curve for train and cv
plot_roc(tr_y, tr_pred, cv_y, cv_pred,2)
```




10 6.7.1 Kaggle Score

```
[21]: # Create a submission file format to submit in Kaggle
temp_id = df_test['id']
eclf_csv = eclf.predict_proba(ts_X)[: ,1]
eclf_df = pd.DataFrame(np.column_stack((temp_id,eclf_csv)),
    ↳columns=['id','target'])
eclf_df['id'] = eclf_df['id'].astype('int32')
eclf_df.to_csv(data_dir+'/submission_eclf.csv', index=False)
```

```
[22]: image = plt.imread(data_dir+'/submission_eclf.png')
plt.figure(figsize=(18,5))
plt.imshow(image)
```

```
[22]: <matplotlib.image.AxesImage at 0x20b16277d48>
```



11 6.8 Voting Classifier (With Stack Classifier + no weights)

```
[23]: # Voting Classifier (See Docs: http://rasbt.github.io/mlxtend/user\_guide/
      ↪ classifier/EnsembleVoteClassifier/)
eclf = EnsembleVoteClassifier(clfs=[clf1, clf2, clf3, clf4, sclf])
# Fit the train data
eclf.fit(tr_X, tr_y)
```

```
[23]: EnsembleVoteClassifier(clfs=[CalibratedClassifierCV(base_estimator=LogisticRegression(C=1000,
    class_weight='balanced',
    dual=False,
    fit_intercept=True,
    intercept_scaling=1,
    l1_ratio=None,
    max_iter=100,
    multi_class='auto',
    n_jobs=None,
    penalty='l1',
    random_state=42,
    solver='liblinear',
    tol=0.0001,
    verbose=0,
    warm_start=False),
                                cv=3, method='sigmoid'),
    CalibratedClass...
    fit_intercept=True,
    intercept_scaling=1,
    l1_ratio=None,
    max_iter=100,
    multi_class='auto',
    n_jobs=None,
    penalty='l1',
    random_state=42,
    solver='liblinear',
    tol=0.0001,
    verbose=0,
    warm_start=False),
    cv=3,
```

```

method='sigmoid'),


store_train_meta_features=False,
use_clones=True,
use_features_in_secondary=False,
use_probas=True, verbose=0)],
refit=True, verbose=0, voting='hard', weights=None)

```

```

[24]: # Predict in probabilities
tr_pred = eclf.predict_proba(tr_X)
cv_pred = eclf.predict_proba(cv_X)
# Plot ROC Curve for train and cv
plot_roc(tr_y, tr_pred, cv_y, cv_pred,2)

```



12 6.8.1 Kaggle Score


```

[25]: # Create a submission file format to submit in Kaggle
temp_id = df_test['id']
eclf_csv = eclf.predict_proba(ts_X)[: ,1]
eclf_df = pd.DataFrame(np.column_stack((temp_id,eclf_csv)),
    columns=['id', 'target'])
eclf_df['id'] = eclf_df['id'].astype('int32')
eclf_df.to_csv(data_dir+'/submission_eclf_stack.csv', index=False)

```

```
[26]: image = plt.imread(data_dir+'/submission_eclf_stack.png')
plt.figure(figsize=(18,5))
plt.imshow(image)
```

```
[26]: <matplotlib.image.AxesImage at 0x20b29c1be88>
```



13 6.9 Voting Classifier (without Stack Classifier + weights)

```
[27]: # Voting Classifier (See Docs: http://rasbt.github.io/mlxtend/user\_guide/classifier/EnsembleVoteClassifier/)
eclf = EnsembleVoteClassifier(clfs=[clf1,clf2,clf3,clf4], weights=[0.3,0.2,0.
    ↪2,0.3])
# Fit the train data
eclf.fit(tr_X,tr_y)
```

```
[27]: EnsembleVoteClassifier(clfs=[CalibratedClassifierCV(base_estimator=LogisticRegression(C=1000,
    class_weight='balanced',
    dual=False,
    fit_intercept=True,
    intercept_scaling=1,
    l1_ratio=None,
    max_iter=100,
    multi_class='auto',
    n_jobs=None,
    penalty='l1',
    random_state=42,
    solver='liblinear',
    tol=0.0001,
    verbose=0,
    warm_start=False),
    cv=3, method='sigmoid'),
    CalibratedClass...
    max_features='auto',
    max_leaf_nodes=None,
    max_samples=None,
    min_impurity_decrease=0.0,
    min_impurity_split=None,
```

```

min_samples_leaf=1,
min_samples_split=2,
min_weight_fraction_leaf=0.0,
n_estimators=300,
n_jobs=None,
oob_score=False,
random_state=None,
verbose=0,
warm_start=False),


                                cv=3, method='sigmoid')],
                                refit=True, verbose=0, voting='hard',
                                weights=[0.3, 0.2, 0.2, 0.3])

```

```

[28]: # Predict in probabilities
tr_pred = eclf.predict_proba(tr_X)
cv_pred = eclf.predict_proba(cv_X)
# Plot ROC Curve for train and cv
plot_roc(tr_y, tr_pred, cv_y, cv_pred,2)

```




14 6.9.1 Kaggle Score

```
[29]: # Create a submission file format to submit in Kaggle
temp_id = df_test['id']
eclf_csv = eclf.predict_proba(ts_X)[: ,1]
eclf_df = pd.DataFrame(np.column_stack((temp_id,eclf_csv)),
    ↪columns=['id', 'target'])
eclf_df['id'] = eclf_df['id'].astype('int32')
eclf_df.to_csv(data_dir+'/submission_eclf_weights.csv', index=False)
```

```
[30]: image = plt.imread(data_dir+'/submission_eclf_weights.png')
plt.figure(figsize=(18,5))
plt.imshow(image)
```

[30]: <matplotlib.image.AxesImage at 0x20b29ebb348>



15 6.10 Voting Classifier (with Stack Classifier + weights)

```
[31]: # Voting Classifier (See Docs: http://rasbt.github.io/mlxtend/user\_guide/
    ↪classifier/EnsembleVoteClassifier/)
eclf = EnsembleVoteClassifier(clfs=[clf1,clf2,clf3,clf4,sclf], weights=[0.3,0.
    ↪1,0.15,0.15,0.3])
# Fit the train data
eclf.fit(tr_X,tr_y)
```

```
[31]: EnsembleVoteClassifier(clfs=[CalibratedClassifierCV(base_estimator=LogisticRegression(C=1000,
    class_weight='balanced',
    dual=False,
    fit_intercept=True,
    intercept_scaling=1,
    l1_ratio=None,
    max_iter=100,
    multi_class='auto',
    n_jobs=None,
    penalty='l1',
    random_state=42,
    solver='liblinear',
```



```

tol=0.0001,
verbose=0,
warm_start=False),

                                cv=3, method='sigmoid'),
                                CalibratedClass...
                                intercept_scaling=1,
                                l1_ratio=None,
                                max_iter=100,
                                multi_class='auto',
                                n_jobs=None,
                                penalty='l1',
                                random_state=42,
                                solver='liblinear',
                                tol=0.0001,
                                verbose=0,
                                warm_start=False),

                                cv=3,
                                method='sigmoid'),


                                store_train_meta_features=False,
                                use_clones=True,
                                use_features_in_secondary=False,
                                use_probabilities=True, verbose=0)],
                                refit=True, verbose=0, voting='hard',
                                weights=[0.3, 0.1, 0.15, 0.15, 0.3])

```

```

[32]: # Predict in probabilities
tr_pred = eclf.predict_proba(tr_X)
cv_pred = eclf.predict_proba(cv_X)
# Plot ROC Curve for train and cv
plot_roc(tr_y, tr_pred, cv_y, cv_pred,2)

```



16 6.10.1 Kaggle Score

```
[33]: # Create a submission file format to submit in Kaggle
temp_id = df_test['id']
eclf_csv = eclf.predict_proba(ts_X)[: ,1]
eclf_df = pd.DataFrame(np.column_stack((temp_id,eclf_csv)),
    ↪columns=['id','target'])
eclf_df['id'] = eclf_df['id'].astype('int32')
eclf_df.to_csv(data_dir+'/submission_eclf_stack_weights.csv', index=False)
```

```
[34]: image = plt.imread(data_dir+'/submission_eclf_stack_weights.png')
plt.figure(figsize=(18,5))
plt.imshow(image)
```

```
[34]: <matplotlib.image.AxesImage at 0x20b2a0ab348>
```



17 7. Summary of all Models

```
[35]: from prettytable import PrettyTable
x = PrettyTable()
x.field_names = ['Model', 'Hyperparameter', 'cv', 'test']
x.add_row(['kNN', r"{'algorithm': 'kd_tree', 'n_neighbors': 35}", 0.70, 0.589])
x.add_row(['Logistic Regression', r"{'C': 1000, 'penalty': 'l1', 'solver': 'liblinear'}", 0.72, 0.803])
x.add_row(['SVC', r"{'C': 10, 'kernel': 'rbf'}", 0.78, 0.697])
x.add_row(['RandomForest', r"{'max_depth': 5, 'n_estimators': 300}", 0.75, 0.732])
x.add_row(['XGBoost', r"{'max_depth': 3, 'n_estimators': 200}", 0.78, 0.758])
x.add_row(['StackClassifier', '-', 0.78, 0.779])
x.add_row(['Voting Classifier(No stacking + no weights)', '-', 0.77, 0.785])
x.add_row(['Voting Classifier(stacking + no weights)', '-', 0.76, 0.784])
x.add_row(['Voting Classifier(no stacking + weights)', '-', 0.76, 0.785])
x.add_row(['Voting Classifier(stacking + weights)', '-', 0.74, 0.785])
print(x)
```

		Model				Hyperparameter	
cv	test						
kNN				{ 'algorithm': 'kd_tree',			
'n_neighbors': 35							
	0.7	0.589		{ 'C': 1000, 'penalty': 'l1',			
Logistic Regression				{ 'C': 1000, 'penalty': 'l1',			
'solver': 'liblinear'}							
	0.72	0.803		{ 'C': 10, 'kernel':			
SVC				{ 'C': 10, 'kernel':			
'rbf'}							
	0.78	0.697		{ 'max_depth': 5,			
RandomForest				{ 'max_depth': 5,			
'n_estimators': 300}							
	0.75	0.732		{ 'max_depth': 3,			
XGBoost				{ 'max_depth': 3,			
'n_estimators': 200}							
	0.78	0.758					
StackClassifier							
	0.78	0.779					
Voting Classifier(No stacking + no weights)							
	0.77	0.785					
Voting Classifier(stacking + no weights)							
	0.76	0.784					
Voting Classifier(no stacking + weights)							
	0.76	0.785					
Voting Classifier(stacking + weights)							
	0.74	0.785					

[]: