

3_4_Models

April 20, 2020

1 SMOTE + Normalization + ML

1. SMOTE → Oversampling technique (called Synthetic Minority Oversampling Technique)
2. No Feature Engineering applied

2 1. Import Necessary Libraries

```
[1]: # For Computational and random seed purpose
import numpy as np
np.random.seed(42)
# To read csv file
import pandas as pd
# To Split data into train and cv data
from sklearn.model_selection import train_test_split
# To compute AUROC score
# For AUROC Score (Ref: https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc\_auc\_score.html)
from sklearn.metrics import roc_curve, auc
# Oversampling technique: SMOTE
from imblearn.over_sampling import SMOTE
# Data is unbalance, we need Calibrated Model to give confidence probabilities
↳ result
from sklearn.calibration import CalibratedClassifierCV
# For Hyperparameter and CV Fold
from sklearn.model_selection import GridSearchCV, StratifiedKFold
# For plot AUROC graph
import matplotlib.pyplot as plt
# For heatmap
import seaborn as sns
# To ignore warnings
import warnings
warnings.filterwarnings('ignore')
# To standardize the data
from sklearn.preprocessing import MinMaxScaler
```

D:\anaconda3\lib\site-packages\tensorflow\python\framework\dtypes.py:516:
FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated;

```

in a future version of numpy, it will be understood as (type, (1,)) /
'(1,)type'.
_np_qint8 = np.dtype(["qint8", np.int8, 1])
D:\anaconda3\lib\site-packages\tensorflow\python\framework\dtypes.py:517:
FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,)) /
'(1,)type'.
_np_quint8 = np.dtype(["quint8", np.uint8, 1])
D:\anaconda3\lib\site-packages\tensorflow\python\framework\dtypes.py:518:
FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,)) /
'(1,)type'.
_np_qint16 = np.dtype(["qint16", np.int16, 1])
D:\anaconda3\lib\site-packages\tensorflow\python\framework\dtypes.py:519:
FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,)) /
'(1,)type'.
_np_quint16 = np.dtype(["quint16", np.uint16, 1])
D:\anaconda3\lib\site-packages\tensorflow\python\framework\dtypes.py:520:
FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,)) /
'(1,)type'.
_np_qint32 = np.dtype(["qint32", np.int32, 1])
D:\anaconda3\lib\site-packages\tensorflow\python\framework\dtypes.py:525:
FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,)) /
'(1,)type'.
_np_resource = np.dtype(["resource", np.ubyte, 1])
D:\anaconda3\lib\site-packages\tensorboard\compat\tensorflow_stub\dtypes.py:541:
FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,)) /
'(1,)type'.
_np_qint8 = np.dtype(["qint8", np.int8, 1])
D:\anaconda3\lib\site-packages\tensorboard\compat\tensorflow_stub\dtypes.py:542:
FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,)) /
'(1,)type'.
_np_quint8 = np.dtype(["quint8", np.uint8, 1])
D:\anaconda3\lib\site-packages\tensorboard\compat\tensorflow_stub\dtypes.py:543:
FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,)) /
'(1,)type'.
_np_qint16 = np.dtype(["qint16", np.int16, 1])
D:\anaconda3\lib\site-packages\tensorboard\compat\tensorflow_stub\dtypes.py:544:
FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,)) /
'(1,)type'.
_np_quint16 = np.dtype(["quint16", np.uint16, 1])

```

```

D:\anaconda3\lib\site-packages\tensorboard\compat\tensorflow_stub\dtypes.py:545:
FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,)) /
'(1,)type'.
_np_qint32 = np.dtype(["qint32", np.int32, 1])
D:\anaconda3\lib\site-packages\tensorboard\compat\tensorflow_stub\dtypes.py:550:
FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated;
in a future version of numpy, it will be understood as (type, (1,)) /
'(1,)type'.
np_resource = np.dtype(["resource", np.ubyte, 1])

```

3 2. Read train data

```

[2]: # Locate parent directory
data_dir = "./"

# Read csv file and display top 5 rows
df_train = pd.read_csv(data_dir+'/train.csv')
df_train.head(5)

```

```

[2]:   id  target      0      1      2      3      4      5      6      7  ...  \
0    0      1.0 -0.098  2.165  0.681 -0.614  1.309 -0.455 -0.236  0.276  ...
1    1      0.0  1.081 -0.973 -0.383  0.326 -0.428  0.317  1.172  0.352  ...
2    2      1.0 -0.523 -0.089 -0.348  0.148 -0.022  0.404 -0.023 -0.172  ...
3    3      1.0  0.067 -0.021  0.392 -1.637 -0.446 -0.725 -1.035  0.834  ...
4    4      1.0  2.347 -0.831  0.511 -0.021  1.225  1.594  0.585  1.509  ...

      290      291      292      293      294      295      296      297      298      299
0  0.867  1.347  0.504 -0.649  0.672 -2.097  1.051 -0.414  1.038 -1.065
1 -0.165 -1.695 -1.257  1.359 -0.808 -1.624 -0.458 -1.099 -0.936  0.973
2  0.013  0.263 -1.222  0.726  1.444 -1.165 -1.544  0.004  0.800 -1.211
3 -0.404  0.640 -0.595 -0.966  0.900  0.467 -0.562 -0.254 -0.533  0.238
4  0.898  0.134  2.415 -0.996 -1.006  1.378  1.246  1.478  0.428  0.253

[5 rows x 302 columns]

```

```

[3]: df_test = pd.read_csv(data_dir+'/test.csv')
df_test.head(5)

```

```

[3]:   id      0      1      2      3      4      5      6      7      8  ...  \
0  250  0.500 -1.033 -1.595  0.309 -0.714  0.502  0.535 -0.129 -0.687  ...
1  251  0.776  0.914 -0.494  1.347 -0.867  0.480  0.578 -0.313  0.203  ...
2  252  1.750  0.509 -0.057  0.835 -0.476  1.428 -0.701 -2.009 -1.378  ...
3  253 -0.556 -1.855 -0.682  0.578  1.592  0.512 -1.419  0.722  0.511  ...
4  254  0.754 -0.245  1.173 -1.623  0.009  0.370  0.781 -1.763 -1.432  ...

      290      291      292      293      294      295      296      297      298      299

```

```

0 -0.088 -2.628 -0.845  2.078 -0.277  2.132  0.609 -0.104  0.312  0.979
1 -0.683 -0.066  0.025  0.606 -0.353 -1.133 -3.138  0.281 -0.625 -0.761
2 -0.094  0.351 -0.607 -0.737 -0.031  0.701  0.976  0.135 -1.327  2.463
3 -0.336 -0.787  0.255 -0.031 -0.836  0.916  2.411  1.053 -1.601 -1.529
4  2.184 -1.090  0.216  1.186 -0.143  0.322 -0.068 -0.156 -1.153  0.825

```

[5 rows x 301 columns]

4 3. Split and Oversampling data

```

[4]: # Take separate for features value
X = df_train.drop(['id','target'], axis=1)
# Take separate for class value
y = df_train['target'].values
# Take test feature value
ts_X = df_test.drop(['id'], axis=1)
# Split the data into train and cv
tr_X, cv_X, tr_y, cv_y = train_test_split(X, y, test_size=0.1, stratify=y,
    ↪random_state=42)
# SMOTE (Ref: https://imbalanced-learn.readthedocs.io/en/stable/generated/
    ↪imblearn.over_sampling.SMOTE.html)
smote = SMOTE()
# Oversampling using SMOTE technique
tr_X, tr_y = smote.fit_sample(tr_X, tr_y)

```

5 4. Normalization

```

[5]: # Fit and transform on train data
stand_vec = MinMaxScaler()
tr_X = stand_vec.fit_transform(tr_X)
pd.DataFrame(tr_X).head(5)

```

```

[5]:
      0      1      2      3      4      5      6  \
0  0.543185  0.385981  0.307849  0.486355  0.593561  0.548421  0.309127
1  0.536840  0.360374  0.582875  0.362963  0.454180  0.333629  0.769649
2  0.429595  0.668411  0.542508  0.541326  0.731114  0.284675  0.213147
3  0.447605  0.656262  0.623445  0.691033  0.192427  0.596488  0.614632
4  0.161891  0.607477  0.423038  0.435673  0.394366  0.523413  0.904745

      7      8      9  ...    290    291    292    293  \
0  0.730249  0.761705  0.344966  ...  0.533780  0.694508  0.788640  0.198537
1  0.242868  0.711077  0.369924  ...  0.575939  0.328375  0.217732  0.584535
2  0.666423  0.082794  0.230753  ...  0.255424  0.507437  0.690333  0.491118
3  0.600585  0.251618  0.124365  ...  0.473276  0.508581  0.471327  0.417346
4  0.400146  0.468786  0.491751  ...  0.171812  0.736651  0.179865  0.376594

```

	294	295	296	297	298	299
0	0.957192	0.640174	0.456048	0.283599	0.700638	0.688088
1	0.793216	0.474931	0.552811	0.373565	0.372348	0.682507
2	0.707056	0.503144	0.449574	0.383006	0.479899	0.557965
3	0.124977	0.607609	0.468654	0.799889	0.434505	0.337426
4	0.765463	0.710463	0.266780	0.552388	0.473817	0.678042

[5 rows x 300 columns]

```
[6]: # Transform on cv data based on mean and std on train data
cv_X = stand_vec.transform(cv_X)
pd.DataFrame(cv_X).head(5)
```

	0	1	2	3	4	5	6	\
0	0.307818	0.205421	0.471764	0.318129	0.479056	0.361121	0.647591	
1	0.826852	0.712150	0.203466	0.288109	0.579111	0.466300	0.687613	
2	0.199959	0.533832	0.425688	0.530604	0.643314	0.238560	0.845346	
3	0.029267	0.095514	0.616514	0.330994	0.263947	0.682866	0.593263	
4	0.468072	0.616449	0.920082	0.324951	0.685019	0.062433	0.472836	

	7	8	9	...	290	291	292	293	\
0	0.461960	0.555577	0.377538	...	0.679309	0.428680	0.579829	0.394984	
1	0.480432	0.245527	0.666244	...	0.841418	0.348207	0.442563	0.704911	
2	0.520666	1.004949	0.696277	...	0.327747	0.746949	0.579465	0.368443	
3	0.445318	0.485535	0.665398	...	0.662904	0.436499	0.851265	0.358830	
4	0.399232	0.474496	0.175127	...	0.397248	0.969489	0.361369	0.505747	

	294	295	296	297	298	299
0	0.286595	0.414154	0.284668	0.681970	0.396084	0.517621
1	0.616724	0.475415	0.696082	0.266753	0.425456	0.538032
2	0.287321	0.778494	0.525554	0.464458	0.295357	0.619678
3	0.739888	0.542157	0.359796	0.350796	0.523661	0.822995
4	0.766008	0.505078	0.490630	0.472233	0.589230	0.328018

[5 rows x 300 columns]

```
[7]: # Transform on test data based on mean and std value from train data
ts_X = stand_vec.transform(ts_X)
pd.DataFrame(ts_X).head(5)
```

	0	1	2	3	4	5	6	\
0	0.576955	0.354766	0.179817	0.520078	0.338760	0.593650	0.635820	
1	0.633442	0.718692	0.404281	0.722417	0.310774	0.589748	0.643607	
2	0.832788	0.642991	0.493374	0.622612	0.382294	0.757893	0.411988	
3	0.360827	0.201121	0.365953	0.572515	0.760563	0.595424	0.281963	
4	0.628940	0.502056	0.744139	0.143470	0.471008	0.570238	0.680369	

	7	8	9	...	290	291	292	293	\
0	0.455925	0.396079	0.772420	...	0.479097	-0.035278	0.348079	0.891745	
1	0.422275	0.565474	0.786168	...	0.374140	0.453280	0.506463	0.584117	
2	0.112107	0.264560	0.534687	...	0.478038	0.532799	0.391407	0.303448	
3	0.611558	0.624096	0.619289	...	0.435350	0.315789	0.548334	0.450993	
4	0.157096	0.254282	0.302665	...	0.879873	0.258009	0.541234	0.705329	

	294	295	296	297	298	299
0	0.545982	0.838949	0.552470	0.494632	0.522623	0.714240
1	0.532197	0.312591	-0.085860	0.565902	0.383623	0.436772
2	0.590604	0.608254	0.614991	0.538874	0.279484	0.950885
3	0.444586	0.642915	0.859455	0.708812	0.238837	0.314304
4	0.570288	0.547155	0.437138	0.485006	0.305296	0.689683

[5 rows x 300 columns]

6 5. Apply ML Models (with hyperparameter)

```
[8]: def hyperparameter_model(models, params):
    """
    Hyperparameter tuning with StratifiedKFold follow by GridSearchCV follow by
    ↳ CalibratedClassifier

    Parameters:
    models: Instance of the model
    params: list of parameters with value fr tuning (dict)

    Return:
    grid_clf: return gridsearch model
    """
    # Perform KCrossValidation with stratified target
    str_cv = StratifiedKFold(n_splits=10, random_state=42)
    # Perform Hyperparamter using GridSearchCV
    grid_clf = GridSearchCV(models, params, cv=str_cv, return_train_score=True,
    ↳ scoring='roc_auc')
    # Fit the train model to evaluate score
    grid_clf.fit(tr_X, tr_y)
    return grid_clf

# Ref: https://scikit-learn.org/stable/auto_examples/model_selection/plot_roc.
↳ html
def plot_roc(try_true, try_pred, cvy_true, cvy_pred, n_classes):
    """
    Compute ROC curve and ROC area for each class
```

```

Parameters:
try_true: train true label
try_pred: train predict probabilities value
cvy_true: cv true label
cvy_pred: cv predict probabilities value
n_classes: number of unique classes

Return:
Plot of ROC Curve for train and cv data
'''

# For train
tr_fpr = dict()
tr_tpr = dict()
tr_roc_auc = dict()
for i in range(n_classes):
    tr_fpr[i], tr_tpr[i], _ = roc_curve(try_true, try_pred[:, i])
    tr_roc_auc[i] = auc(tr_fpr[i], tr_tpr[i])

# For cv
cv_fpr = dict()
cv_tpr = dict()
cv_roc_auc = dict()
for i in range(n_classes):
    cv_fpr[i], cv_tpr[i], _ = roc_curve(cvy_true, cvy_pred[:, i])
    cv_roc_auc[i] = auc(cv_fpr[i], cv_tpr[i])

# Line thickness
lw = 2

# Plot roc for train
plt.plot(tr_fpr[1], tr_tpr[1], color='red',
         lw=lw, label='ROC curve for Train (area = %0.2f)' % tr_roc_auc[1])

# Plot roc for cv
plt.plot(cv_fpr[1], cv_tpr[1], color='green',
         lw=lw, label='ROC curve for CV (area = %0.2f)' % cv_roc_auc[1])

plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic: train vs cv')
plt.legend(loc="lower right")
plt.show()

def plot_feature_importance(model, model_name, top_n = 10):
    '''
    Plot the feature importance on the basis of model.

```

```

Parameters:
model: Instance of model
model_name: Name of the model
top_n: Number of feature you want to print top features

Return:
df: DataFrame that return feature names with coefficient in descending order
Plot the feature importance
'''

# Numpy Column Stack (See Docs: https://docs.scipy.org/doc/numpy-1.10.1/reference/generated/numpy.column\_stack.html)

column_name = df_train.drop(['id', 'target'], axis=1).columns
if model_name == 'log_model':
    feat_imp_coef = model.coef_.ravel()
else:
    feat_imp_coef = model.feature_importances_
temp = pd.DataFrame(data=np.column_stack((column_name, feat_imp_coef)),
→columns=['col_name', 'coef'])
temp = temp.sort_values(by='coef', ascending=False).reset_index()
df = temp
temp = temp[:top_n]
plt.figure(figsize=(20,5))
sns.barplot(data=temp, y='coef', x='col_name', order=temp['col_name'])
plt.grid()
plt.show()
return df

```

7 5.1 kNN

```

[13]: # Import KNN
from sklearn.neighbors import KNeighborsClassifier

```

```

[14]: # kNN (See Docs: https://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsClassifier.html)

# List of params
params = {'n_neighbors':np.arange(3,51,2).tolist(), 'algorithm': ['kd_tree',
→'brute']}
# Instance of knn model
knn_model = KNeighborsClassifier()
# Call hyperparameter for find the best params as possible
knn_clf = hyperparameter_model(knn_model, params)


```




```
[15]: cv_pvt = pd.pivot_table(pd.DataFrame(knn_clf.cv_results_),
    ↪ values='mean_test_score', index='param_n_neighbors', \
        columns='param_algorithm')
tr_pvt = pd.pivot_table(pd.DataFrame(knn_clf.cv_results_),
    ↪ values='mean_train_score', index='param_n_neighbors', \
        columns='param_algorithm')

plt.title('Train Hyperparameter')
sns.heatmap(tr_pvt, annot=True)
plt.show()

plt.title('CV Hyperparameter')
sns.heatmap(cv_pvt, annot=True)
plt.show()
```






```
[16]: print(knn_clf.best_params_)
print('CV Score',knn_clf.score(cv_X,cv_y))
```

```
{'algorithm': 'kd_tree', 'n_neighbors': 47}
CV Score 0.75
```

```
[17]: clf = CalibratedClassifierCV(knn_clf, cv=3)
clf.fit(tr_X,tr_y)

tr_pred = clf.predict_proba(tr_X)
cv_pred = clf.predict_proba(cv_X)

# Plot ROC cureve of train and cv data
plot_roc(tr_y, tr_pred, cv_y, cv_pred, 2)
```




8 5.1.1 Kaggle Score

```
[18]: # Create a submission format to make submission in Kaggle
temp_id = df_test['id']
knn_csv = clf.predict_proba(ts_X)[: ,1]
knn_df = pd.DataFrame(np.column_stack((temp_id,knn_csv)),
    ↪columns=['id', 'target'])
knn_df['id'] = knn_df['id'].astype('int32')
knn_df.to_csv(data_dir+'/submission_knn.csv', index=False)
```

```
[19]: image = plt.imread(data_dir+'/submission_knn.png')
plt.figure(figsize=(18,5))
plt.imshow(image)
```

```
[19]: <matplotlib.image.AxesImage at 0x1e60b674e48>
```



8.1 5.2 Logistic Regression

```
[9]: # Import Logistic Regression
from sklearn.linear_model import LogisticRegression
```


```
[21]: # LogisticRegression (See Docs: https://scikit-learn.org/stable/modules/generated/sklearn.linear\_model.LogisticRegression.html)


# List of hyperparameter that has to be tuned
params = {'penalty':['l1', 'l2', 'elasticnet'], 'C':[10**i for i in
→range(-4,5)], 'solver':['liblinear','sag']}
# Instance of logistic regression
log_model = LogisticRegression(random_state=42, class_weight='balanced')
# Call hyperparameter to find the best params
log_clf = hyperparameter_model(log_model, params)
```

```
[22]: cv_pvt = pd.pivot_table(pd.DataFrame(log_clf.cv_results_),
→values='mean_test_score', index='param_penalty', \
        columns='param_C')
tr_pvt = pd.pivot_table(pd.DataFrame(log_clf.cv_results_),
→values='mean_train_score', index='param_penalty', \
        columns='param_C')

plt.title('Train Hyperparameter')
sns.heatmap(tr_pvt, annot=True)
plt.show()

plt.title('CV Hyperparameter')
sns.heatmap(cv_pvt, annot=True)
plt.show()
```






```
[23]: print(log_clf.best_params_)
      print('cv score',log_clf.score(cv_X,cv_y))

{'C': 1000, 'penalty': 'l1', 'solver': 'liblinear'}
cv score 0.7013888888888889
```

```
[24]: clf = CalibratedClassifierCV(log_clf, cv=3)
      clf.fit(tr_X,tr_y)

      tr_pred = clf.predict_proba(tr_X)
      cv_pred = clf.predict_proba(cv_X)


      # Plot ROC curve of train and cv data
      plot_roc(tr_y, tr_pred, cv_y, cv_pred, 2)
```



```
[25]: log_model = LogisticRegression(**log_clf.best_params_, random_state=42,
    ↪class_weight='balanced')
log_model.fit(tr_X, tr_y)
```

```
[25]: LogisticRegression(C=1000, class_weight='balanced', dual=False,
    fit_intercept=True, intercept_scaling=1, l1_ratio=None,
    max_iter=100, multi_class='auto', n_jobs=None, penalty='l1',
    random_state=42, solver='liblinear', tol=0.0001, verbose=0,
    warm_start=False)
```

```
[26]: df = plot_feature_importance(log_model, 'log_model', 20)
```




8.2 5.2.1 Kaggle Score

```
[27]: # Create a submission format to make submission in Kaggle
temp_id = df_test['id']
log_csv = clf.predict_proba(ts_X)[: ,1]
log_df = pd.DataFrame(np.column_stack((temp_id,log_csv)),
    ↪columns=['id', 'target'])
log_df['id'] = log_df['id'].astype('int32')
log_df.to_csv(data_dir+'/submission_log.csv', index=False)
```

```
[28]: image = plt.imread(data_dir+'/submission_log.png')
plt.figure(figsize=(18,5))
plt.imshow(image)
```

```
[28]: <matplotlib.image.AxesImage at 0x1e602177508>
```



8.3 5.3 SVC

```
[10]: # Import SVC
from sklearn.svm import SVC
```

```
[30]: # SVC (See Docs: https://scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html)
    ↪SVC.html)


# List of hyperparameter that has to be tuned
params = {'C':[10**i for i in range(-4,5)], 'kernel':
    ↪['linear', 'poly', 'sigmoid', 'rbf']}
# Instance of SVC
svc_model = SVC(class_weight='balanced', random_state=42, probability=True)
# Call hyperparameter to find the best parameters
svc_clf = hyperparameter_model(svc_model, params)
```


```
[31]: cv_pvt = pd.pivot_table(pd.DataFrame(svc_clf.cv_results_),
    ↪values='mean_test_score', index='param_C', \
    columns='param_kernel')
tr_pvt = pd.pivot_table(pd.DataFrame(svc_clf.cv_results_),
    ↪values='mean_train_score', index='param_C', \
    columns='param_kernel')
```



```
plt.title('Train Hyperparameter')
sns.heatmap(tr_pvt, annot=True)
plt.show()

plt.title('CV Hyperparameter')
sns.heatmap(cv_pvt, annot=True)
plt.show()
```





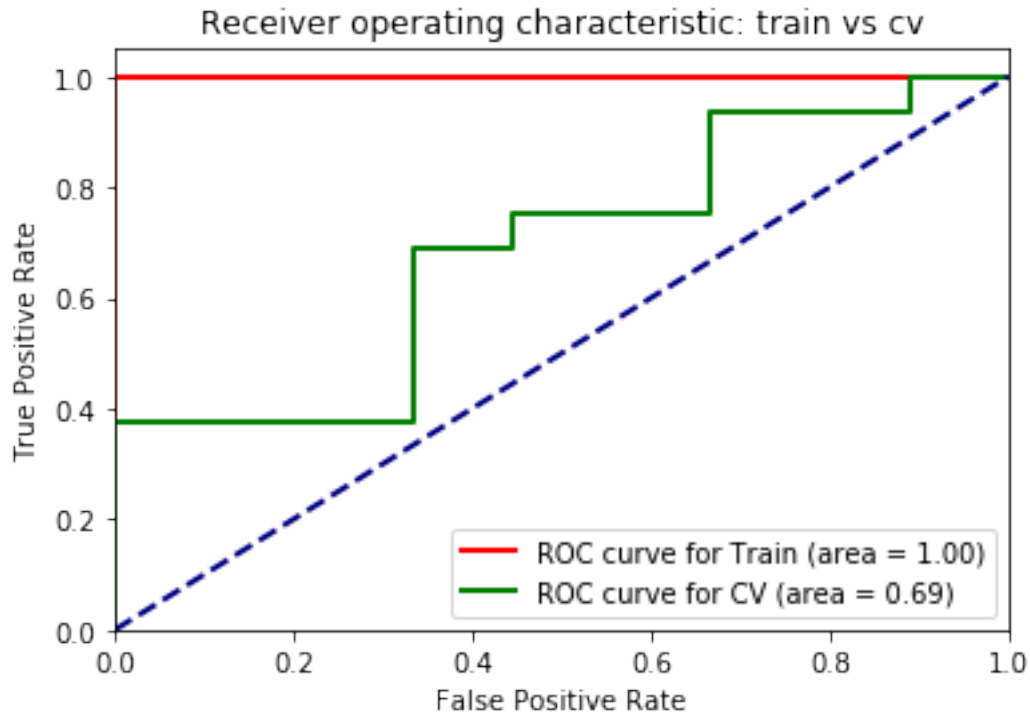
```
[32]: print(svc_clf.best_params_)
      print('cv Score',svc_clf.score(cv_X,cv_y))
```

```
{'C': 10, 'kernel': 'rbf'}
cv Score 0.6458333333333334
```

```
[33]: clf = CalibratedClassifierCV(svc_clf, cv=3)
      clf.fit(tr_X,tr_y)

      tr_pred = clf.predict_proba(tr_X)
      cv_pred = clf.predict_proba(cv_X)

      # Plot ROC curve of this model
      plot_roc(tr_y, tr_pred, cv_y, cv_pred, 2)
```



8.4 5.3.1 Kaggle Score

```
[34]: # Create a submssion format to make submission in Kaggle
temp_id = df_test['id']
svc_csv = clf.predict_proba(ts_X)[: ,1]
svc_df = pd.DataFrame(np.column_stack((temp_id,svc_csv)),
    ↪columns=['id', 'target'])
svc_df['id'] = svc_df['id'].astype('int32')
svc_df.to_csv(data_dir+'/submission_svc.csv', index=False)
```

```
[35]: image = plt.imread(data_dir+'/submission_svc.png')
plt.figure(figsize=(18,5))
plt.imshow(image)
```

```
[35]: <matplotlib.image.AxesImage at 0x1e602300188>
```



9 5.4 RandomForest



```
[11]: # Import Random Forest
      from sklearn.ensemble import RandomForestClassifier

[37]: # RandomForest (See Docs: https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestClassifier.html)

      # List of hyperparameter that has to be tuned
      params = {'n_estimators': [10, 20, 30, 40, 50, 100, 200, 300, 400], 'max_depth': [2, 3, 5, 7]}
      # Instance of RandomForest
      rf_model = RandomForestClassifier(random_state=42)
      # Perform GridSearchCV to find best parameters
      rf_clf = hyperparameter_model(rf_model, params)

[38]: # Ref: https://stackoverflow.com/questions/48791709/how-to-plot-a-heat-map-on-pivot-table-after-grid-search

      # Plotting of hyperparameter of train and cv score
      pvt_tr = pd.pivot_table(pd.DataFrame(rf_clf.cv_results_),
                              values='mean_train_score', index='param_n_estimators',
                              columns='param_max_depth')
      pvt_cv = pd.pivot_table(pd.DataFrame(rf_clf.cv_results_),
                              values='mean_test_score', index='param_n_estimators',
                              columns='param_max_depth')
      plt.figure(1)
      plt.title('Hyperparameter for Train')
      sns.heatmap(pvt_tr, annot=True)
      plt.figure(2)
      plt.title('Hyperparameter for CV')
      sns.heatmap(pvt_cv, annot=True)
      plt.show()
```



```
[39]: print(rf_clf.best_params_)
```

```
{'max_depth': 5, 'n_estimators': 300}
```

```
[40]: # Calibrate the model
```

```
clf = CalibratedClassifierCV(rf_clf, cv=3)
clf.fit(tr_X, tr_y)
```

```
[40]: CalibratedClassifierCV(base_estimator=GridSearchCV(cv=StratifiedKFold(n_splits=1
0, random_state=42, shuffle=False),
```

```
error_score=nan,
```

```
estimator=RandomForestClassifier(bootstrap=True,
```

```
ccp_alpha=0.0,
```

```
class_weight=None,
```

```
criterion='gini',
```

```
max_depth=None,
```

```
max_features='auto',
```

```
max_leaf_nodes=None,
```

```
max_samples=None,
```

```
min_impurity_decrease=0.0,
```

```
min_impurity_split=None,
```

```
mi...
```

```
min_samples_split=2,
```

```
min_weight_fraction_leaf=0.0,
```

```
n_estimators=100,
```

```
n_jobs=None,
```

```
oob_score=False,
```

```
random_state=42,
```

```
verbose=0,
```

```
warm_start=False),
```

```
iid='deprecated',
```

```
n_jobs=None,
```

```
param_grid={'max_depth': [2,
3,
5,
7],
```

```
'n_estimators':
```

```
[10,
```

```
20,
```

```
30,
```

```
40,
```

```
50,
```

```
100,
```

```
200,
```

```
300,
```

```
400]}],
```

```
pre_dispatch='2*n_jobs',
```

```

refit=True,
return_train_score=True,
scoring='roc_auc',
verbose=0),


cv=3, method='sigmoid')

```

```

[41]: # Plot ROC Curve of train and cv
plot_roc(tr_y, clf.predict_proba(tr_X), cv_y, clf.predict_proba(cv_X), 2)

```




```

[42]: rf_clf = RandomForestClassifier(**rf_clf.best_params_, random_state=42)
rf_clf.fit(tr_X, tr_y)

# Plot the feature importance on the basis of this model
df = plot_feature_importance(rf_clf, 'rf', 30)

```




9.1 5.4.1 Kaggle Score

```
[43]: temp_id = df_test['id']
      rf_csv = clf.predict_proba(ts_X)[: ,1]
      rf_df = pd.DataFrame(np.column_stack((temp_id,rf_csv)), columns=['id','target'])
      rf_df['id'] = rf_df['id'].astype('int32')
      rf_df.to_csv(data_dir+'/submission_rf.csv', index=False)
```

```
[44]: image = plt.imread(data_dir+'/submission_rf.png')
      plt.figure(figsize=(18,5))
      plt.imshow(image)
```

```
[44]: <matplotlib.image.AxesImage at 0x1e601f14d88>
```



9.2 5.5 Xgboost

```
[12]: # Import Xgboost
      from xgboost import XGBClassifier
```

```
[46]: # Xgboost (See Docs: https://xgboost.readthedocs.io/en/latest/python/python\_api.html)

      # List of hyperparameter that has to be tuned
      params = {'max_depth':[2,3,5,7], 'n_estimators':[10,20,50,100,200,300,400]}
      # Instance of XGBoost Model
      xgb_model = XGBClassifier(scale_pos_weight=0.5)
      # Call hyperparameter to find the best parameters
      xgb_clf = hyperparameter_model(xgb_model, params)
```

```
[47]: # Ref: https://stackoverflow.com/questions/48791709/how-to-plot-a-heat-map-on-pivot-table-after-grid-search


      # Plotting of hyperparameter of train and cv score
      pvt_tr = pd.pivot_table(pd.DataFrame(xgb_clf.cv_results_),
                              values='mean_train_score', index='param_n_estimators',
                              columns='param_max_depth')
```




```

pvt_cv = pd.pivot_table(pd.DataFrame(xgb_clf.cv_results_),
    ↪ values='mean_test_score', index='param_n_estimators',
    ↪ columns='param_max_depth')
plt.figure(1)
plt.title('Hyperparameter for Train')
sns.heatmap(pvt_tr, annot=True)
plt.figure(2)
plt.title('Hyperparameter for CV')
sns.heatmap(pvt_cv, annot=True)
plt.show()

```






```
[48]: print(xgb_clf.best_params_)
      print('cv Score',xgb_clf.score(cv_X,cv_y))
```

```
{'max_depth': 3, 'n_estimators': 200}
cv Score 0.7847222222222223
```


```
[49]: # Instance of randomforest with best parameters
xgb_clf = XGBClassifier(**xgb_clf.best_params_, random_state=42,
    ↪scale_pos_weight=0.5)
# Fit the model
xgb_clf.fit(tr_X,tr_y)
# Calibrate the model
clf = CalibratedClassifierCV(xgb_clf, cv=3)
clf.fit(tr_X, tr_y)

tr_pred = clf.predict_proba(tr_X)
cv_pred = clf.predict_proba(cv_X)

# Plot ROC curve of train and cv
plot_roc(tr_y, tr_pred, cv_y, cv_pred, 2)
```



```
[50]: # Instance of XGBoost model with best parameters
df = plot_feature_importance(xgb_clf, 'xgb', 10)
```




9.3 5.5.1 Kaggle Score

```
[51]: temp_id = df_test['id']
xgb_csv = clf.predict_proba(ts_X)[: , 1]
xgb_df = pd.DataFrame(np.column_stack((temp_id, xgb_csv)),
    columns=['id', 'target'])
xgb_df['id'] = xgb_df['id'].astype('int32')
xgb_df.to_csv(data_dir + '/submission_xgb.csv', index=False)
```

```
[52]: image = plt.imread(data_dir+'/submission_xgb.png')
plt.figure(figsize=(18,5))
plt.imshow(image)
```

```
[52]: <matplotlib.image.AxesImage at 0x1e602efe4c8>
```



9.4 5.6 Stacking Model

```
[13]: # Import Stacking Classifier
from mlxtend.classifier import StackingClassifier
```

```
[14]: # StackClassifier (See Docs: http://rasbt.github.io/mlxtend/user\_guide/classifier/StackingClassifier/#methods)

# Classifier 1: Logistic Regression with best params
clf1 = LogisticRegression(C = 1000, penalty = 'l1', solver = 'liblinear',
    ↪class_weight='balanced', random_state=42)
clf1.fit(tr_X, tr_y)
clf1 = CalibratedClassifierCV(clf1, cv=3)

# Classifier 2: SVC with best params
clf2 = SVC(C=10, kernel='rbf', random_state=42, class_weight='balanced',
    ↪probability=True)
clf2.fit(tr_X, tr_y)
clf2 = CalibratedClassifierCV(clf2, cv=3)

# Classifier 3: XGBoost with best params
clf3 = XGBClassifier(max_depth=3, n_estimators=200, scale_pos_weight=0.5)
clf3.fit(tr_X, tr_y)
clf3 = CalibratedClassifierCV(clf3, cv=3)

# Classifier 4: RF with best params
clf4 = RandomForestClassifier(max_depth=5, n_estimators=300)
clf4.fit(tr_X, tr_y)
clf4 = CalibratedClassifierCV(clf4, cv=3)

# Stack Classifier
sclf = StackingClassifier(classifiers=[clf1, clf2, clf3, clf4],
    ↪meta_classifier=clf1, use_proba=True)
```


```
# Fit the model
sclf.fit(tr_X, tr_y)

# Predict in probabilities
tr_pred = sclf.predict_proba(tr_X)
cv_pred = sclf.predict_proba(cv_X)
```

```
[55]: # Score after stacking classifier
sclf.score(cv_X, cv_y)
```

```
[55]: 0.68
```

```
[56]: # Plot ROC Curve for train and cv
plot_roc(tr_y, tr_pred, cv_y, cv_pred,2)
```




9.5 5.6.1 Kaggle Score

```
[57]: temp_id = df_test['id']
sclf_csv = sclf.predict_proba(ts_X)[: ,1]
sclf_df = pd.DataFrame(np.column_stack((temp_id,sclf_csv)),
    ↪columns=['id', 'target'])
sclf_df['id'] = sclf_df['id'].astype('int32')
```

```
sclf_df.to_csv(data_dir+'/submission_sclf.csv', index=False)
```

```
[58]: image = plt.imread(data_dir+'/submission_sclf.png')  
plt.figure(figsize=(18,5))  
plt.imshow(image)
```


```
[58]: <matplotlib.image.AxesImage at 0x1e602cec848>
```



10 5.7 Voting Classifier (Without Stack Classifier + no weights)

```
[15]: # Import Voting Classifier  
from mlxtend.classifier import EnsembleVoteClassifier
```

```
[16]: # Voting Classifier (See Docs: http://rasbt.github.io/mlxtend/user\_guide/classifier/EnsembleVoteClassifier/)  
ecf = EnsembleVoteClassifier(clfs=[clf1, clf2, clf3, clf4])  
# Fit the train data  
ecf.fit(tr_X, tr_y)  
  
# Predict in probabilities  
tr_pred = ecf.predict_proba(tr_X)  
cv_pred = ecf.predict_proba(cv_X)  
# Plot ROC Curve for train and cv  
plot_roc(tr_y, tr_pred, cv_y, cv_pred, 2)
```




11 5.7.1 Kaggle Score

```
[17]: # Create a submission file format to submit in Kaggle
temp_id = df_test['id']
eclf_csv = eclf.predict_proba(ts_X)[: ,1]
eclf_df = pd.DataFrame(np.column_stack((temp_id,eclf_csv)),
    ↪columns=['id','target'])
eclf_df['id'] = eclf_df['id'].astype('int32')
eclf_df.to_csv(data_dir+'/submission_eclf.csv', index=False)
```

```
[20]: image = plt.imread(data_dir+'/submission_eclf.png')
plt.figure(figsize=(18,5))
plt.imshow(image)
```


```
[20]: <matplotlib.image.AxesImage at 0x1355f974a48>
```



12 5.8 Voting Classifier (With Stack Classifier + no weights)

```
[21]: # Voting Classifier (See Docs: http://rasbt.github.io/mlxtend/user\_guide/classifier/EnsembleVoteClassifier/)
eclf = EnsembleVoteClassifier(clfs=[clf1, clf2, clf3, clf4, sclf])
# Fit the train data
eclf.fit(tr_X, tr_y)

# Predict in probabilities
tr_pred = eclf.predict_proba(tr_X)
cv_pred = eclf.predict_proba(cv_X)
# Plot ROC Curve for train and cv
plot_roc(tr_y, tr_pred, cv_y, cv_pred, 2)
```



13 5.8.1 Kaggle Score


```
[19]: # Create a submission file format to submit in Kaggle
temp_id = df_test['id']
eclf_csv = eclf.predict_proba(ts_X)[: , 1]
eclf_df = pd.DataFrame(np.column_stack((temp_id, eclf_csv)),
    columns=['id', 'target'])
eclf_df['id'] = eclf_df['id'].astype('int32')
```



```
eclf_df.to_csv(data_dir+'/submission_eclf_stack.csv', index=False)
```


```
[22]: image = plt.imread(data_dir+'/submission_eclf_stack.png')  
plt.figure(figsize=(18,5))  
plt.imshow(image)
```

```
[22]: <matplotlib.image.AxesImage at 0x1355fa5e448>
```



14 5.9 Voting Classifier (without Stack Classifier + weights)

```
[23]: # Voting Classifier (See Docs: http://rasbt.github.io/mlxtend/user\_guide/  
      ↪ classifier/EnsembleVoteClassifier/)  
eclf = EnsembleVoteClassifier(clfs=[clf1,clf2,clf3,clf4], weights=[0.3,0.1,0.  
      ↪ 3,0.3])  
# Fit the train data  
eclf.fit(tr_X,tr_y)  
  
# Predict in probabilities  
tr_pred = eclf.predict_proba(tr_X)  
cv_pred = eclf.predict_proba(cv_X)  
# Plot ROC Curve for train and cv  
plot_roc(tr_y, tr_pred, cv_y, cv_pred,2)
```




15 5.9.1 Kaggle Score

```
[24]: # Create a submission file format to submit in Kaggle
temp_id = df_test['id']
eclf_csv = eclf.predict_proba(ts_X)[: ,1]
eclf_df = pd.DataFrame(np.column_stack((temp_id,eclf_csv)),
    ↪columns=['id','target'])
eclf_df['id'] = eclf_df['id'].astype('int32')
eclf_df.to_csv(data_dir+'/submission_eclf_weights.csv', index=False)
```

```
[25]: image = plt.imread(data_dir+'/submission_eclf_weights.png')
plt.figure(figsize=(18,5))
plt.imshow(image)
```


```
[25]: <matplotlib.image.AxesImage at 0x1355f8a0c48>
```



16 5.10 Voting Classifier (with Stack Classifier + weights)

```
[26]: # Voting Classifier (See Docs: http://rasbt.github.io/mlxtend/user\_guide/classifier/EnsembleVoteClassifier/)
      eclf = EnsembleVoteClassifier(clfs=[clf1,clf2,clf3,clf4,sclf], weights=[0.3,0.
      ↪05,0.15,0.2,0.3])
      # Fit the train data
      eclf.fit(tr_X,tr_y)

      # Predict in probabilities
      tr_pred = eclf.predict_proba(tr_X)
      cv_pred = eclf.predict_proba(cv_X)
      # Plot ROC Curve for train and cv
      plot_roc(tr_y, tr_pred, cv_y, cv_pred,2)
```



17 5.10.1 Kaggle Score

```
[27]: # Create a submission file format to submit in Kaggle
      temp_id = df_test['id']
      eclf_csv = eclf.predict_proba(ts_X)[: ,1]
      eclf_df = pd.DataFrame(np.column_stack((temp_id,eclf_csv)),
      ↪columns=['id', 'target'])
```

```
eclf_df['id'] = eclf_df['id'].astype('int32')
eclf_df.to_csv(data_dir+'/submission_eclf_stack_weights.csv', index=False)
```

```
[ ]: image = plt.imread(data_dir+'/submission_eclf_stack_weights.png')
plt.figure(figsize=(18,5))
plt.imshow(image)
```

18 6. Summary of all Models

```
[29]: from prettytable import PrettyTable
x = PrettyTable()
x.field_names = (['Model', 'Hyperparameter', 'CV score', 'Test score'])
x.add_row(['kNN', r"{'algorithm': 'kd_tree', 'n_neighbors': 47}", 0.69, 0.586])
x.add_row(['Logistic Regression', r"{'C': 1000, 'penalty': 'l1', 'solver': 'liblinear'}", 0.65, 0.772])
x.add_row(['SVC', r"{'C': 10, 'kernel': 'rbf'}", 0.69, 0.695])
x.add_row(['RandomForest', r"{'max_depth': 5, 'n_estimators': 300}", 0.74, 0.736])
x.add_row(['XGBoost', r"{'max_depth': 3, 'n_estimators': 200}", 0.84, 0.751])
x.add_row(['Stack Classifier', '-', 0.72, 0.768])
x.add_row(['Voting Classifier(No stacking + no weights)', '-', 0.74, 0.775])
x.add_row(['Voting Classifier(stacking + no weights)', '-', 0.73, 0.777])
x.add_row(['Voting Classifier(no stacking + weights)', '-', 0.75, 0.779])
x.add_row(['Voting Classifier(stacking + weights)', '-', 0.71, 0.776])
print(x)
```

Model	Hyperparameter	CV score	Test score
kNN	{'algorithm': 'kd_tree', 'n_neighbors': 47}	0.69	0.586
Logistic Regression	{'C': 1000, 'penalty': 'l1', 'solver': 'liblinear'}	0.65	0.772
SVC	{'C': 10, 'kernel': 'rbf'}	0.69	0.695
RandomForest	{'max_depth': 5, 'n_estimators': 300}	0.74	0.736
XGBoost	{'max_depth': 3, 'n_estimators': 200}	0.84	0.751
Stack Classifier	-	0.72	0.768
Voting Classifier(No stacking + no weights)	-	0.74	0.775
Voting Classifier(stacking + no weights)	-	0.73	0.777
Voting Classifier(no stacking + weights)	-	0.75	0.779
Voting Classifier(stacking + weights)	-	0.71	0.776

	0.73		0.777		
	Voting Classifier(no stacking + weights)				
	0.75		0.779		-
	Voting Classifier(stacking + weights)				
	0.71		0.776		-
+-----+					
-----+-----+					

[]: