

## 4\_1\_Models

April 20, 2020

### 1 Standardization + ML Classification Model with/without Top Features

1. No oversampling techniques applied
2. No feature engineering applied

#### 2 1. Import Necessary Libraries

```
[17]: # For Computational and random seed purpose
import numpy as np
np.random.seed(42)
# To read csv file
import pandas as pd
# To Split data into train and cv data
from sklearn.model_selection import train_test_split
# To compute AUROC score
# For AUROC Score (Ref: https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc\_auc\_score.html)
from sklearn.metrics import roc_curve, auc
# For Hyperparameter and CV Fold
from sklearn.model_selection import GridSearchCV, RepeatedStratifiedKFold, \
    cross_val_score
# For plot AUROC graph
import matplotlib.pyplot as plt
# Data is unbalance, we need Calibrated Model to give confidence probabilities
from sklearn.calibration import CalibratedClassifierCV
# For heatmap
import seaborn as sns
# To ignore warnings
import warnings
warnings.filterwarnings('ignore')
# To standardize the data
from sklearn.preprocessing import StandardScaler
import tqdm
```

## 3 2. Read train data

```
[2]: # Locate parent directory
data_dir = "./"

# Read csv file and display top 5 rows
df_train = pd.read_csv(data_dir+'/train.csv')
df_train.head(5)
```

```
[2]:   id  target      0      1      2      3      4      5      6      7  ... \
0    0      1.0 -0.098  2.165  0.681 -0.614  1.309 -0.455 -0.236  0.276  ...
1    1      0.0  1.081 -0.973 -0.383  0.326 -0.428  0.317  1.172  0.352  ...
2    2      1.0 -0.523 -0.089 -0.348  0.148 -0.022  0.404 -0.023 -0.172  ...
3    3      1.0  0.067 -0.021  0.392 -1.637 -0.446 -0.725 -1.035  0.834  ...
4    4      1.0  2.347 -0.831  0.511 -0.021  1.225  1.594  0.585  1.509  ...

      290      291      292      293      294      295      296      297      298      299
0  0.867  1.347  0.504 -0.649  0.672 -2.097  1.051 -0.414  1.038 -1.065
1 -0.165 -1.695 -1.257  1.359 -0.808 -1.624 -0.458 -1.099 -0.936  0.973
2  0.013  0.263 -1.222  0.726  1.444 -1.165 -1.544  0.004  0.800 -1.211
3 -0.404  0.640 -0.595 -0.966  0.900  0.467 -0.562 -0.254 -0.533  0.238
4  0.898  0.134  2.415 -0.996 -1.006  1.378  1.246  1.478  0.428  0.253
```

[5 rows x 302 columns]

```
[3]: df_test = pd.read_csv(data_dir+'/test.csv')
df_test.head(5)
```

```
[3]:   id      0      1      2      3      4      5      6      7      8  ... \
0  250  0.500 -1.033 -1.595  0.309 -0.714  0.502  0.535 -0.129 -0.687  ...
1  251  0.776  0.914 -0.494  1.347 -0.867  0.480  0.578 -0.313  0.203  ...
2  252  1.750  0.509 -0.057  0.835 -0.476  1.428 -0.701 -2.009 -1.378  ...
3  253 -0.556 -1.855 -0.682  0.578  1.592  0.512 -1.419  0.722  0.511  ...
4  254  0.754 -0.245  1.173 -1.623  0.009  0.370  0.781 -1.763 -1.432  ...

      290      291      292      293      294      295      296      297      298      299
0 -0.088 -2.628 -0.845  2.078 -0.277  2.132  0.609 -0.104  0.312  0.979
1 -0.683 -0.066  0.025  0.606 -0.353 -1.133 -3.138  0.281 -0.625 -0.761
2 -0.094  0.351 -0.607 -0.737 -0.031  0.701  0.976  0.135 -1.327  2.463
3 -0.336 -0.787  0.255 -0.031 -0.836  0.916  2.411  1.053 -1.601 -1.529
4  2.184 -1.090  0.216  1.186 -0.143  0.322 -0.068 -0.156 -1.153  0.825
```

[5 rows x 301 columns]

## 4 3. Take train and test values from DataFrame

```
[4]: # Take separate for features value
tr_X = df_train.drop(['id', 'target'], axis=1)
# Take separate for class value
tr_y = df_train['target'].values
# Take test feature value
ts_X = df_test.drop(['id'], axis=1)
```

Note: Don't worry about splitting train data into train and cv. I apply Stratify CV technique while modelling

## 5 4. Standardization

```
[5]: stand_vec = StandardScaler()
tr_X = stand_vec.fit_transform(tr_X)
pd.DataFrame(tr_X).head(5)
```

```
[5]:
```

	0	1	2	3	4	5	6	\
0	-0.121736	2.176002	0.503692	-0.609972	1.265232	-0.469388	-0.266814	
1	1.061577	-0.939278	-0.539790	0.320974	-0.415729	0.340017	1.134681	
2	-0.548290	-0.061678	-0.505465	0.144689	-0.022827	0.431232	-0.054798	
3	0.043868	0.005829	0.220265	-1.623118	-0.433148	-0.752470	-1.062122	
4	2.332208	-0.798306	0.336970	-0.022683	1.183942	1.678890	0.550393	

	7	8	9	...	290	291	292	293	\
0	0.210682	-2.296917	1.758518	...	0.814697	1.257605	0.509875	-0.664341	
1	0.291718	0.042547	-0.320787	...	-0.207701	-1.876475	-1.339295	1.430147	
2	-0.267006	0.180835	0.144993	...	-0.031358	0.140793	-1.302542	0.769883	
3	0.805660	0.561388	0.234415	...	-0.444478	0.529204	-0.644150	-0.994994	
4	1.525391	0.025911	2.125050	...	0.845409	0.007888	2.516556	-1.026286	

	294	295	296	297	298	299
0	0.699064	-1.921131	1.040182	-0.421724	1.022411	-0.965720
1	-0.703385	-1.464704	-0.522377	-1.108796	-0.939658	1.136883
2	1.430611	-1.021786	-1.646923	-0.002459	0.785850	-1.116348
3	0.915116	0.553031	-0.630069	-0.261240	-0.539094	0.378584
4	-0.891009	1.432111	1.242103	1.476000	0.416098	0.394060

[5 rows x 300 columns]

```
[6]: ts_X = stand_vec.transform(ts_X)
pd.DataFrame(ts_X).head(5)
```

```
[6]:
```

	0	1	2	3	4	5	6	\
0	0.478452	-0.998843	-1.728417	0.304138	-0.692502	0.533981	0.500624	
1	0.755461	0.934060	-0.648649	1.332140	-0.840565	0.510915	0.543426	

2	1.733024	0.531992	-0.220077	0.825072	-0.462180	1.504847	-0.729665	
3	-0.581411	-1.814892	-0.833024	0.570547	1.539102	0.544465	-1.444348	
4	0.733381	-0.216549	0.986204	-1.609253	0.007173	0.395585	0.745488	
	7	8	9	...	290	291	292	293 \
0	-0.221157	-0.675928	1.233779	...	-0.131418	-2.837717	-0.906667	2.180115
1	-0.417350	0.249460	1.297652	...	-0.720882	-0.198166	0.006893	0.644714
2	-2.225742	-1.394404	0.129271	...	-0.137362	0.231456	-0.656750	-0.756131
3	0.686238	0.569706	0.522334	...	-0.377110	-0.940990	0.248408	-0.019722
4	-1.963439	-1.450551	-0.948706	...	2.119444	-1.253162	0.207456	1.249696
	294	295	296	297	298	299		
0	-0.200209	2.159692	0.582494	-0.110786	0.300799	1.143073		
1	-0.272227	-0.990908	-3.297498	0.275379	-0.630538	-0.652084		
2	0.032901	0.778832	0.962520	0.128937	-1.328295	2.674115		
3	-0.729917	0.986299	2.448452	1.049714	-1.600639	-1.444429		
4	-0.073231	0.413112	-0.118535	-0.162943	-1.155347	0.984192		

[5 rows x 300 columns]

## 6 5. Apply ML Models (with hyperparameter)

```
[8]: def hyperparameter_model(models, params):
    """
    Hyperparameter tuning with StratifiedKFold follow by GridSearchCV follow by
    CalibratedClassifier

    Parameters:
    models: Instance of the model
    params: list of parameters with value fr tuning (dict)

    Return:
    grid_clf: return gridsearch model
    """
    # Random shuffle after every iteration with stratify
    str_cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=5, random_state=42)
    # Find the right hyperparameter for the model
    grid_clf = GridSearchCV(models, params, cv=str_cv, return_train_score=True,
    scoring='roc_auc')
    # Fit on train data
    grid_clf.fit(tr_X, tr_y)
    return grid_clf

def plot_feature_importance(model, model_name, top_n = 10):
    """
    Compute ROC curve and ROC area for each class
```

```

Parameters:
try_true: train true label
try_pred: train predict probabilities value
cvy_true: cv true label
cvy_pred: cv predict probabilities value
n_classes: number of unique classes

Return:
Plot of ROC Curve for train and cv data
'''

column_name = df_train.drop(['id', 'target'], axis=1).columns
if model_name == 'log_model':
    feat_imp_coef = model.coef_.ravel()
else:
    feat_imp_coef = model.feature_importances_
temp = pd.DataFrame(data=np.column_stack((column_name, feat_imp_coef)),
→columns=['col_name', 'coef'])
temp = temp.sort_values(by='coef', ascending=False).reset_index()
df = temp
temp = temp[:top_n]
plt.figure(figsize=(20,5))
sns.barplot(data=temp, y='coef', x='col_name', order=temp['col_name'])
plt.grid()
plt.show()
return df

```

```

[12]: def forward_selection_model(model, top_n=10):
    top_column = []
    exist_score = 0
    for n in range(top_n):
        print('for {} feature'.format(n+1))
        flag = 0
        for i in tqdm.tqdm_notebook(range(tr_X.shape[1])):
            # Finding for first top feature
            if len(top_column) == 0:
                str_cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=5,
→random_state=42)
                score = cross_val_score(model, tr_X[:
→, [i]], tr_y, cv=str_cv, scoring='roc_auc')
                if exist_score < np.mean(score):
                    top_current = i
                    exist_score = np.mean(score)
                    flag = 1
            # Excluded apart from top column
            elif i not in top_column:

```

```

        str_cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=5,
        ↪random_state=42)
        score = cross_val_score(model, tr_X[:, np.
        ↪concatenate((top_column, [i]))], tr_y, cv=str_cv, scoring='roc_auc')
        if exist_score < np.mean(score):
            top_current = i
            exist_score = np.mean(score)
            flag = 1

    if flag == 1:
        print('Current top feature {} and score: {}'.
        ↪format(top_current, exist_score))
        print('Appended to top column')
        top_column.append(top_current)
        print(top_column)
    else:
        break
    return top_column

```

## 7 5.1 kNN


```


[10]: # Import KNN
      from sklearn.neighbors import KNeighborsClassifier

[11]: # kNN (See Docs: 

```

```
plt.title('CV Hyperparameter')
sns.heatmap(cv_pvt, annot=True)
plt.show()
```





```
[15]: print(knn_clf.best_params_)
```

```
{'algorithm': 'kd_tree', 'n_neighbors': 35}
```

```
[18]: clf = CalibratedClassifierCV(knn_clf, cv=3)
      clf.fit(tr_X, tr_y)
```

```
[18]: CalibratedClassifierCV(base_estimator=GridSearchCV(cv=RepeatedStratifiedKfold(n_
repeats=5, n_splits=10, random_state=42),
                                error_score=nan,
                                estimator=KNeighborsClassifier(algorithm='auto',
leaf_size=30,
metric='minkowski',
metric_params=None,
n_jobs=None,
n_neighbors=5,
p=2,
weights='uniform'),
                                iid='deprecated',
                                n_jobs=None,
                                param_grid={'algorithm':
['kd_tree',
'brute']},
```



```

'n_neighbors':
[3,
5,
7,
9,
11,
13,
15,
17,
19,
21,
23,
25,
27,
29,
31,
33,
35,
37,
39,
41,
43,
45,
47,
49]],

pre_dispatch='2*n_jobs',
refit=True,
return_train_score=True,
scoring='roc_auc',
verbose=0),

cv=3, method='sigmoid')

```

### 8 5.1.1 Kaggle Score without top features

```

[19]: # Create a submission format to make submission in Kaggle
temp_id = df_test['id']
knn_csv = clf.predict_proba(ts_X)[: ,1]
knn_df = pd.DataFrame(np.column_stack((temp_id,knn_csv)),
    ↪columns=['id', 'target'])
knn_df['id'] = knn_df['id'].astype('int32')
knn_df.to_csv(data_dir+'/submission_knn.csv', index=False)


```

```

[21]: image = plt.imread(data_dir+'/submission_knn.png')
plt.figure(figsize=(18,5))
plt.imshow(image)

```

[21]: <matplotlib.image.AxesImage at 0x266dc241548>



## 8.1 5.2 Logistic Regression

```
[22]: # Import Logistic Regression
from sklearn.linear_model import LogisticRegression
```


```
[23]: # LogisticRegression (See Docs: https://scikit-learn.org/stable/modules/generated/sklearn.linear\_model.LogisticRegression.html)


# List of hyperparameter that has to be tuned
params = {'penalty':['l1', 'l2', 'elasticnet'], 'C':[10**i for i in
    range(-4,5)], 'solver':['liblinear', 'sag']}
# Instance of Logsitic Regression
log_model = LogisticRegression(random_state=42, class_weight='balanced')
# Call hyperparameter to get the best parameters of this model
log_clf = hyperparameter_model(log_model, params)
```

```
[24]: cv_pvt = pd.pivot_table(pd.DataFrame(log_clf.cv_results_),
    values='mean_test_score', index='param_penalty', \
        columns='param_C')
tr_pvt = pd.pivot_table(pd.DataFrame(log_clf.cv_results_),
    values='mean_train_score', index='param_penalty', \
        columns='param_C')

plt.title('Train Hyperparameter')
sns.heatmap(tr_pvt, annot=True)
plt.show()

plt.title('CV Hyperparameter')
sns.heatmap(cv_pvt, annot=True)
plt.show()
```






```
[25]: print(log_clf.best_params_)
```

```
{'C': 0.1, 'penalty': 'l1', 'solver': 'liblinear'}
```

```
[26]: log_model = LogisticRegression(**log_clf.best_params_, class_weight='balanced',
    ↪ random_state=42)
log_model.fit(tr_X, tr_y)
```

```
[26]: LogisticRegression(C=0.1, class_weight='balanced', dual=False,
    fit_intercept=True, intercept_scaling=1, l1_ratio=None,
    max_iter=100, multi_class='auto', n_jobs=None, penalty='l1',
    random_state=42, solver='liblinear', tol=0.0001, verbose=0,
    warm_start=False)
```

```
[27]: df = plot_feature_importance(log_model, 'log_model', 20)
```



```
[28]: clf = CalibratedClassifierCV(log_clf, cv=3)
      clf.fit(tr_X, tr_y)
```

```
[28]: CalibratedClassifierCV(base_estimator=GridSearchCV(cv=RepeatedStratifiedKFold(n_
repeats=5, n_splits=10, random_state=42),

error_score=nan,

estimator=LogisticRegression(C=1.0,
class_weight='balanced',
dual=False,
fit_intercept=True,
intercept_scaling=1,
l1_ratio=None,
max_iter=100,
multi_class='auto',
n_jobs=None,
penalty='l2',
random_state=42,
solver='lbfgs',
tol=0.0001,
verbose=0,
warm_start=False),

iid='deprecated',
n_jobs=None,
param_grid={'C': [0.0001,
                  0.001,
                  0.01, 0.1,
                  1, 10, 100,
                  1000,
                  10000],
            'penalty': ['l1',
                       'l2'],

'elasticnet'],

'solver':

['liblinear',
'sag']},
```

```

pre_dispatch='2*n_jobs',
refit=True,
return_train_score=True,
scoring='roc_auc',
verbose=0),

cv=3, method='sigmoid')

```

## 9 5.2.1 Kaggle Score without top features

```

[29]: # Create a submission format to make submission in Kaggle
temp_id = df_test['id']
log_csv = clf.predict_proba(ts_X)[: ,1]
log_df = pd.DataFrame(np.column_stack((temp_id,log_csv)),
    ↪columns=['id', 'target'])
log_df['id'] = log_df['id'].astype('int32')
log_df.to_csv(data_dir+'/submission_log.csv', index=False)

```

```


[30]: image = plt.imread(data_dir+'/submission_log.png')
plt.figure(figsize=(18,5))
plt.imshow(image)

```

```

[30]: <matplotlib.image.AxesImage at 0x266dedfdd08>

```



## 10 5.2.2 Kaggle Score Using top 10 features based on Logistic Regression Model

```

[31]: top_column = forward_selection_model(log_model)

```

for 1 feature

```

HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))

```

Current top feature 33 and score: 0.7347222222222222

Appended to top column

[33]

for 2 feature

```
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))
```

Current top feature 65 and score: 0.7876388888888887

Appended to top column

[33, 65]

for 3 feature

```
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))
```

Current top feature 217 and score: 0.8005555555555557

Appended to top column

[33, 65, 217]

for 4 feature

```
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))
```

Current top feature 117 and score: 0.8156944444444443

Appended to top column

[33, 65, 217, 117]

for 5 feature

```
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))
```

Current top feature 91 and score: 0.8325

Appended to top column

[33, 65, 217, 117, 91]

for 6 feature

```
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))
```

Current top feature 73 and score: 0.8423611111111111

Appended to top column

[33, 65, 217, 117, 91, 73]

for 7 feature

```
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))
```

Current top feature 16 and score: 0.8493055555555555

Appended to top column

[33, 65, 217, 117, 91, 73, 16]

for 8 feature

```
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))
```

Current top feature 226 and score: 0.8527777777777778

Appended to top column

[33, 65, 217, 117, 91, 73, 16, 226]

for 9 feature

```
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))
```

Current top feature 82 and score: 0.8559722222222222

Appended to top column

[33, 65, 217, 117, 91, 73, 16, 226, 82]

for 10 feature

```
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))
```

Current top feature 90 and score: 0.8590277777777778

Appended to top column

[33, 65, 217, 117, 91, 73, 16, 226, 82, 90]


```
[37]: # Store the top column name into set
log_top10_feat = set(top_column)
# Fit Logistic Regression on top features only
log_model_top = LogisticRegression(**log_clf.best_params_,
    ↪class_weight='balanced', random_state=42)
log_model_top.fit(tr_X[:,top_column],tr_y)
# Calibrate it in top features only
clf = CalibratedClassifierCV(log_model_top, cv=3)
clf.fit(tr_X[:,top_column],tr_y)
# Predict the probabilities of 1
ts_pred = clf.predict_proba(ts_X[:,top_column])[:,1]
```

```
[34]: # Create a submission format to make submission in Kaggle
temp_id = df_test['id']
df_tspred = pd.DataFrame(np.column_stack((temp_id,ts_pred)),
    ↪columns=['id','target'])
df_tspred['id'] = df_tspred['id'].astype('int32')
df_tspred.to_csv(data_dir+'/submission_log_top10.csv', index=False)
```

```
[36]: image = plt.imread(data_dir+'/log_csv_top10.png')
plt.figure(figsize=(18,5))
plt.imshow(image)
```

```
[36]: <matplotlib.image.AxesImage at 0x266e864dd08>
```





## 11 5.2.3 Kaggle Score Using top 20 features based on Logistic Regression Model

```
[39]: top_column = forward_selection_model(log_model,20)
```

```
for 1 feature
```

```
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))
```

```
Current top feature 33 and score: 0.7347222222222222
```

```
Appended to top column
```

```
[33]
```

```
for 2 feature
```

```
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))
```

```
Current top feature 65 and score: 0.7876388888888887
```

```
Appended to top column
```

```
[33, 65]
```

```
for 3 feature
```

```
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))
```

```
Current top feature 217 and score: 0.8005555555555557
```

```
Appended to top column
```

```
[33, 65, 217]
```

```
for 4 feature
```

```
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))
```

```
Current top feature 117 and score: 0.8156944444444443
```

```
Appended to top column
```

```
[33, 65, 217, 117]
```

```
for 5 feature
```

```
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))
```

Current top feature 91 and score: 0.8325

Appended to top column

[33, 65, 217, 117, 91]

for 6 feature

```
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))
```

Current top feature 73 and score: 0.8423611111111111

Appended to top column

[33, 65, 217, 117, 91, 73]

for 7 feature

```
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))
```

Current top feature 16 and score: 0.8493055555555555

Appended to top column

[33, 65, 217, 117, 91, 73, 16]

for 8 feature

```
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))
```

Current top feature 226 and score: 0.8527777777777778

Appended to top column

[33, 65, 217, 117, 91, 73, 16, 226]

for 9 feature

```
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))
```

Current top feature 82 and score: 0.8559722222222222

Appended to top column

[33, 65, 217, 117, 91, 73, 16, 226, 82]

for 10 feature

```
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))
```

Current top feature 90 and score: 0.8590277777777778

Appended to top column

[33, 65, 217, 117, 91, 73, 16, 226, 82, 90]

for 11 feature

```
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))
```

Current top feature 45 and score: 0.8591666666666669

Appended to top column

[33, 65, 217, 117, 91, 73, 16, 226, 82, 90, 45]

for 12 feature


```
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))
```

```
[40]: # Store the top column name into set
log_top20_feat = set(top_column)
# Fit Logistic Regression on top features only
log_model_top = LogisticRegression(**log_clf.best_params_,
    ↪class_weight='balanced', random_state=42)
log_model_top.fit(tr_X[:,top_column],tr_y)
# Calibrate it in top features only
clf = CalibratedClassifierCV(log_model_top, cv=3)
clf.fit(tr_X[:,top_column],tr_y)
# Predict the probabilities of 1
ts_pred = clf.predict_proba(ts_X[:,top_column])[:,1]
```

```
[41]: # Create a submission format to make submission in Kaggle
temp_id = df_test['id']
df_tspred = pd.DataFrame(np.column_stack((temp_id,ts_pred)),
    ↪columns=['id', 'target'])
df_tspred['id'] = df_tspred['id'].astype('int32')
df_tspred.to_csv(data_dir+'/submission_log_top20.csv', index=False)
```

```
[42]: image = plt.imread(data_dir+'/log_csv_top20.png')
plt.figure(figsize=(18,5))
plt.imshow(image)
```

```
[42]: <matplotlib.image.AxesImage at 0x266e9ff0b48>
```



## 12 5.3 SVC

```
[43]: # Import SVC
      from sklearn.svm import SVC



[44]: # SVC (See Docs: https://scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html)

      # List of hyperparameter that has to be tuned
      params = {'C': [10**i for i in range(-4, 5)], 'kernel':
        ↳ ['linear', 'poly', 'sigmoid', 'rbf']}
      # Instance of SVC
      svc_model = SVC(class_weight='balanced', random_state=42, probability=True)
      # Call hyperparameter to find the best parameters
      svc_clf = hyperparameter_model(svc_model, params)

[45]: cv_pvt = pd.pivot_table(pd.DataFrame(svc_clf.cv_results_),
        ↳ values='mean_test_score', index='param_C', \
          columns='param_kernel')
      tr_pvt = pd.pivot_table(pd.DataFrame(svc_clf.cv_results_),
        ↳ values='mean_train_score', index='param_C', \
          columns='param_kernel')

      plt.title('Train Hyperparameter')
      sns.heatmap(tr_pvt, annot=True)
      plt.show()

      plt.title('CV Hyperparameter')
      sns.heatmap(cv_pvt, annot=True)
      plt.show()
```



```
[46]: print(svc_clf.best_params_)
```

```
{'C': 10, 'kernel': 'rbf'}
```

```
[47]: svc_model = SVC(**svc_clf.best_params_, class_weight='balanced',  
    ↪random_state=42, probability=True)  
svc_model.fit(tr_X, tr_y)  
  
clf = CalibratedClassifierCV(svc_clf, cv=3)  
clf.fit(tr_X, tr_y)
```

```
[47]: CalibratedClassifierCV(base_estimator=GridSearchCV(cv=RepeatedStratifiedKFold(n_  
repeats=5, n_splits=10, random_state=42),  
  
error_score=nan,  
estimator=SVC(C=1.0,  
  
break_ties=False,  
  
cache_size=200,  
  
class_weight='balanced',  
  
coef0=0.0,  
  
decision_function_shape='ovr',  
  
degree=3,  
gamma='scale',  
kernel='rbf',  
max_iter=-1,  
  
probability=True,  
random_state=42,  
  
shrinking=True,  
tol=0.001,  
verbose=False),  
iid='deprecated',  
n_jobs=None,  
param_grid={'C': [0.0001,  
0.001,  
0.01, 0.1,  
1, 10, 100,  
1000,  
10000],  
'kernel':  
  
['linear',  
'poly',  
'sigmoid',  
'rbf']}},  
  
pre_dispatch='2*n_jobs',  
refit=True,  
return_train_score=True,  
scoring='roc_auc',  
verbose=0),
```


```
cv=3, method='sigmoid')
```

### 13 5.3.1 Kaggle Score without top features

```
[48]: # Create a submission format to make submission in Kaggle
temp_id = df_test['id']
svc_csv = clf.predict_proba(ts_X)[: ,1]
svc_df = pd.DataFrame(np.column_stack((temp_id,svc_csv)),
    ↪columns=['id', 'target'])
svc_df['id'] = svc_df['id'].astype('int32')
svc_df.to_csv(data_dir+'/submission_svc.csv', index=False)
```

```
[49]: image = plt.imread(data_dir+'/submission_svc.png')
plt.figure(figsize=(18,5))
plt.imshow(image)
```

```
[49]: <matplotlib.image.AxesImage at 0x266e9bc10c8>
```



### 14 5.3.2 Kaggle Score Using top 10 features based on SVC Model

```
[50]: top_column = forward_selection_model(svc_model)
```

```
for 1 feature
```

```
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))
```

```
Current top feature 33 and score: 0.7088888888888889
```

```
Appended to top column
```

```
[33]
```

```
for 2 feature
```

```
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))
```

```
Current top feature 167 and score: 0.7579166666666667
```

```
Appended to top column
```

```

[33, 167]
for 3 feature
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))

Current top feature 199 and score: 0.7611111111111112
Appended to top column
[33, 167, 199]
for 4 feature
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))

Current top feature 65 and score: 0.7808333333333333
Appended to top column
[33, 167, 199, 65]
for 5 feature
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))

Current top feature 224 and score: 0.8034722222222223
Appended to top column
[33, 167, 199, 65, 224]
for 6 feature
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))

Current top feature 1 and score: 0.8180555555555555
Appended to top column
[33, 167, 199, 65, 224, 1]
for 7 feature
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))

Current top feature 208 and score: 0.8381944444444445
Appended to top column
[33, 167, 199, 65, 224, 1, 208]
for 8 feature
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))

Current top feature 201 and score: 0.867638888888889
Appended to top column

```



```

[33, 167, 199, 65, 224, 1, 208, 201]
for 9 feature

HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))

Current top feature 253 and score: 0.8869444444444444
Appended to top column
[33, 167, 199, 65, 224, 1, 208, 201, 253]
for 10 feature

HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))

```

```

[51]: # Store the top column name into set
svc_top10_feat = set(top_column)
# Fit SVC on top features only
svc_model_top = SVC(**svc_clf.best_params_, class_weight='balanced',
    ↪random_state=42, probability=True)
svc_model_top.fit(tr_X[:,top_column],tr_y)
# Calibrate it in top features only
clf = CalibratedClassifierCV(svc_model_top, cv=3)
clf.fit(tr_X[:,top_column],tr_y)
# Predict the probabilities of 1
ts_pred = clf.predict_proba(ts_X[:,top_column])[:,1]

```

```

[52]: # Create a submission format to make submission in Kaggle
temp_id = df_test['id']
df_tspred = pd.DataFrame(np.column_stack((temp_id,ts_pred)),
    ↪columns=['id','target'])
df_tspred['id'] = df_tspred['id'].astype('int32')
df_tspred.to_csv(data_dir+'/submission_svc_top10.csv', index=False)

```

```


[54]: image = plt.imread(data_dir+'/svc_csv_top10.png')
plt.figure(figsize=(18,5))
plt.imshow(image)

```

```

[54]: <matplotlib.image.AxesImage at 0x266e8512f88>

```



## 15 5.4 RandomForest



```
[55]: # Import Random Forest
      from sklearn.ensemble import RandomForestClassifier

[56]: # RandomForest (See Docs: https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestClassifier.html)

      # List of hyperparameter that has to be tuned
      params = {'n_estimators': [10, 20, 30, 40, 50, 100, 200, 300, 400], 'max_depth': [2, 3, 5, 7]}
      # Instance of RandomForest
      rf_model = RandomForestClassifier(random_state=42)
      # Perform GridSearchCV to find best parameters
      rf_clf = hyperparameter_model(rf_model, params)

[57]: # Ref: https://stackoverflow.com/questions/48791709/how-to-plot-a-heat-map-on-pivot-table-after-grid-search

      # Plotting of hyperparameter of train and cv score
      pvt_tr = pd.pivot_table(pd.DataFrame(rf_clf.cv_results_),
                              values='mean_train_score', index='param_n_estimators',
                              columns='param_max_depth')
      pvt_cv = pd.pivot_table(pd.DataFrame(rf_clf.cv_results_),
                              values='mean_test_score', index='param_n_estimators',
                              columns='param_max_depth')
      plt.figure(1)
      plt.title('Hyperparameter for Train')
      sns.heatmap(pvt_tr, annot=True)
      plt.figure(2)
      plt.title('Hyperparameter for CV')
      sns.heatmap(pvt_cv, annot=True)
      plt.show()
```



```
[58]: print(rf_clf.best_params_)
```

```
{'max_depth': 2, 'n_estimators': 300}
```

```
[59]: # Instance of randomforest with best parameters
rf_model = RandomForestClassifier(**rf_clf.best_params_, random_state=42)
# Fit the model
rf_model.fit(tr_X, tr_y)
# Calibrate the model
clf = CalibratedClassifierCV(rf_clf, cv=3)
clf.fit(tr_X, tr_y)
```

```
[59]: CalibratedClassifierCV(base_estimator=GridSearchCV(cv=RepeatedStratifiedKFold(n_
repeats=5, n_splits=10, random_state=42),
```

```
error_score=nan,
```

```
estimator=RandomForestClassifier(bootstrap=True,
```

```
ccp_alpha=0.0,
```

```
class_weight=None,
```

```
criterion='gini',
```

```
max_depth=None,
```

```
max_features='auto',
```

```
max_leaf_nodes=None,
```

```
max_samples=None,
```

```
min_impurity_decrease=0.0,
```

```
min_impurity_split=None,
```

```
min_samples_split=2,
```

```
min_weight_fraction_leaf=0.0,
```

```
n_estimators=100,
```

```
n_jobs=None,
```

```
oob_score=False,
```

```
random_state=42,
```

```
verbose=0,
```

```
warm_start=False),
```

```
iid='deprecated',
```

```
n_jobs=None,
```

```
param_grid={'max_depth': [2,
3,
5,
7],
```

```
'n_estimators':
```

```
[10,
```

```
20,
```

```
30,
```

```
40,
```

```
50,
```

```
100,
```


```
200,
```

```
300,
400]],
```

```
pre_dispatch='2*n_jobs',
refit=True,
return_train_score=True,
scoring='roc_auc',
verbose=0),

cv=3, method='sigmoid')
```

```
[60]: # Plot the feature importance based on this model
df = plot_feature_importance(rf_model, 'rf',30)
```




## 16 5.4.1 Kaggle Score without top features

```
[61]: # Create a submission file format to submit in kaggle
temp_id = df_test['id']
rf_csv = clf.predict_proba(ts_X)[: ,1]
rf_df = pd.DataFrame(np.column_stack((temp_id,rf_csv)), columns=['id','target'])
rf_df['id'] = rf_df['id'].astype('int32')
rf_df.to_csv(data_dir+'/submission_rf.csv', index=False)
```

```
[62]: image = plt.imread(data_dir+'/submission_rf.png')
plt.figure(figsize=(18,5))
plt.imshow(image)
```

```
[62]: <matplotlib.image.AxesImage at 0x266e9c5f348>
```



## 17 5.4.2 Kaggle Score Using top 10 features based on Random-Forest Model

```
[63]: top_column = forward_selection_model(rf_model)

for 1 feature
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))

Current top feature 33 and score: 0.7097222222222223
Appended to top column
[33]
for 2 feature
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))

Current top feature 65 and score: 0.765
Appended to top column
[33, 65]
for 3 feature
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))

Current top feature 91 and score: 0.7863888888888888
Appended to top column
[33, 65, 91]
for 4 feature
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))

Current top feature 199 and score: 0.7888888888888889
Appended to top column
[33, 65, 91, 199]
for 5 feature
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))

Current top feature 217 and score: 0.8051388888888889
Appended to top column
[33, 65, 91, 199, 217]
for 6 feature
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))
```

```

Current top feature 214 and score: 0.8094444444444444
Appended to top column
[33, 65, 91, 199, 217, 214]
for 7 feature
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))

```

```

Current top feature 189 and score: 0.8162500000000001
Appended to top column
[33, 65, 91, 199, 217, 214, 189]
for 8 feature
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))

```

```

Current top feature 132 and score: 0.8290277777777777
Appended to top column
[33, 65, 91, 199, 217, 214, 189, 132]
for 9 feature
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))

```

```

[64]: # Store the top column name into set
rf_top10_feat = set(top_column)

```

```

[65]: # Fit RF on top features only
rf_model_top = RandomForestClassifier(**rf_clf.best_params_, random_state=42)
rf_model_top.fit(tr_X[:,top_column],tr_y)
# Calibrate it in top features only
clf = CalibratedClassifierCV(rf_model_top, cv=3)
clf.fit(tr_X[:,top_column],tr_y)
# Predict the probabilities of 1
ts_pred = clf.predict_proba(ts_X[:,top_column])[:,1]

```


```

[66]: # Create a submssion format to make submission in Kaggle
temp_id = df_test['id']
df_tspred = pd.DataFrame(np.column_stack((temp_id,ts_pred)),
    ↪columns=['id','target'])
df_tspred['id'] = df_tspred['id'].astype('int32')
df_tspred.to_csv(data_dir+'/submission_rf_top10.csv', index=False)

```

```
[67]: image = plt.imread(data_dir+'/rf_csv_top10.png')
plt.figure(figsize=(18,5))
plt.imshow(image)
```

```
[67]: <matplotlib.image.AxesImage at 0x266e9e12a08>
```



## 18 5.5 XGBoost

```
[69]: # Import Xgboost
from xgboost import XGBClassifier
```



```
[70]: # Xgboost (See Docs: https://xgboost.readthedocs.io/en/latest/python/python\_api.html)
```

```
# List of hyperparameter that has to be tuned
params = {'max_depth': [2, 3, 5, 7], 'n_estimators': [10, 20, 50, 100, 200, 300, 400]}
# Instance of XGBoost Model
xgb_model = XGBClassifier(scale_pos_weight=0.5)
# Call hyperparameter to find the best parameters
xgb_clf = hyperparameter_model(xgb_model, params)
```

```
[71]: # Ref: https://stackoverflow.com/questions/48791709/how-to-plot-a-heat-map-on-pivot-table-after-grid-search
```

```
# Plotting of hyperparameter of train and cv score
pvt_tr = pd.pivot_table(pd.DataFrame(xgb_clf.cv_results_),
    values='mean_train_score', index='param_n_estimators',
    columns='param_max_depth')
pvt_cv = pd.pivot_table(pd.DataFrame(xgb_clf.cv_results_),
    values='mean_test_score', index='param_n_estimators',
    columns='param_max_depth')
plt.figure(1)
plt.title('Hyperparameter for Train')
sns.heatmap(pvt_tr, annot=True)
plt.figure(2)
plt.title('Hyperparameter for CV')
sns.heatmap(pvt_cv, annot=True)
plt.show()
```






```
[72]: print(xgb_clf.best_params_)
```

```
{'max_depth': 2, 'n_estimators': 300}
```

```
[74]: # Instance of randomforest with best parameters
xgb_model = XGBClassifier(**xgb_clf.best_params_, random_state=42,
    ↪ scale_pos_weight=0.5)
# Fit the model
xgb_model.fit(tr_X, tr_y)

# Instance of XGBoost model with best parameters
df = plot_feature_importance(xgb_model, 'xgb', 10)
```



```
[75]: # Calibrate the model
clf = CalibratedClassifierCV(xgb_clf, cv=3)
clf.fit(tr_X, tr_y)
```

```
[75]: CalibratedClassifierCV(base_estimator=GridSearchCV(cv=RepeatedStratifiedKFold(n_
repeats=5, n_splits=10, random_state=42),

error_score=nan,

estimator=XGBClassifier(base_score=None,
booster=None,
colsample_bylevel=None,
colsample_bynode=None,
colsample_bytree=None,
gamma=None,
gpu_id=None,
importance_type='gain',
interaction_constraints=None,
learning_rate=None,

random_state=None,
reg_alpha=None,
reg_lambda=None,
scale_pos_weight=0.5,
subsample=None,
```

m...

```

tree_method=None,
validate_parameters=False,
verbosity=None),

                                iid='deprecated',
                                n_jobs=None,
                                param_grid={'max_depth': [2,
                                                            3,
                                                            5,
                                                            7],
                                'n_estimators':

[10,
20,
50,
100,
200,
300,
400]},

                                pre_dispatch='2*n_jobs',
                                refit=True,
                                return_train_score=True,
                                scoring='roc_auc',
                                verbose=0),

                                cv=3, method='sigmoid')

```

## 19 5.5.1 Kaggle Score without top features

```

[76]: # Create submission file format to submit in Kaggle
temp_id = df_test['id']
xgb_csv = clf.predict_proba(ts_X)[: ,1]
xgb_df = pd.DataFrame(np.column_stack((temp_id,xgb_csv)),
    ↪columns=['id', 'target'])
xgb_df['id'] = xgb_df['id'].astype('int32')
xgb_df.to_csv(data_dir+'/submission_xgb.csv', index=False)

```

```


[77]: image = plt.imread(data_dir+'/submission_xgb.png')
plt.figure(figsize=(18,5))
plt.imshow(image)

```

```

[77]: <matplotlib.image.AxesImage at 0x266f1803908>

```



## 20 5.5.2 Using top 10 features of this model

```
[78]: top_column = forward_selection_model(xgb_model)

for 1 feature
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))

Current top feature 91 and score: 0.6425
Appended to top column
[91]
for 2 feature
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))

Current top feature 120 and score: 0.7095833333333332
Appended to top column
[91, 120]
for 3 feature
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))

Current top feature 33 and score: 0.7473611111111111
Appended to top column
[91, 120, 33]
for 4 feature
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))

Current top feature 108 and score: 0.7788888888888887
Appended to top column
[91, 120, 33, 108]
for 5 feature
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))

Current top feature 141 and score: 0.8011111111111111
Appended to top column
[91, 120, 33, 108, 141]
for 6 feature
```

```
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))
```

Current top feature 246 and score: 0.8295833333333335

Appended to top column

[91, 120, 33, 108, 141, 246]

for 7 feature

```
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))
```

Current top feature 146 and score: 0.8491666666666668

Appended to top column

[91, 120, 33, 108, 141, 246, 146]

for 8 feature

```
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))
```

Current top feature 95 and score: 0.8688888888888889

Appended to top column

[91, 120, 33, 108, 141, 246, 146, 95]

for 9 feature

```
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))
```

Current top feature 230 and score: 0.8731944444444444

Appended to top column

[91, 120, 33, 108, 141, 246, 146, 95, 230]

for 10 feature

```
HBox(children=(FloatProgress(value=0.0, max=300.0), HTML(value='')))
```

```
[79]: # Store the top column name into set
xgb_top10_feat = set(top_column)
```


```
[80]: # Fit RF on top features only
xgb_model_top = XGBClassifier(**xgb_clf.best_params_, random_state=42,
    ↪scale_pos_weight=0.5)
xgb_model_top.fit(tr_X[:,top_column],tr_y)
# Calibrate it in top features only
clf = CalibratedClassifierCV(xgb_model_top, cv=3)
clf.fit(tr_X[:,top_column],tr_y)
# Predict the probabilities of 1
```

```
ts_pred = clf.predict_proba(ts_X[:,top_column])[:,1]
```

```
[81]: # Create a submission format to make submission in Kaggle
temp_id = df_test['id']
df_tspred = pd.DataFrame(np.column_stack((temp_id,ts_pred)),
    ↪columns=['id','target'])
df_tspred['id'] = df_tspred['id'].astype('int32')
df_tspred.to_csv(data_dir+'/submission_xgb_top10.csv', index=False)
```

```
[85]: image = plt.imread(data_dir+'/xgb_csv_top10.png')
plt.figure(figsize=(18,5))
plt.imshow(image)
```

```
[85]: <matplotlib.image.AxesImage at 0x266eaa06a88>
```



## 21 5.6 Stacking Classifier

```
[92]: # Combined all top features of all models
# Ref: https://www.geeksforgeeks.org/union-function-python/
comb_top_feat = list(log_top20_feat.union(svc_top10_feat, rf_top10_feat,
    ↪xgb_top10_feat))
```

```
[86]: # Import Stacking Classifier
from mlxtend.classifier import StackingClassifier
```

```
[93]: # StackClassifier (See Docs: http://rasbt.github.io/mlxtend/user_guide/
    ↪classifier/StackingClassifier/#methods)

# Classifier 1: Logistic Regression with best params
clf1 = LogisticRegression(C = 0.1, penalty = 'l1', solver = 'liblinear',
    ↪class_weight='balanced', random_state=42)
clf1.fit(tr_X,tr_y)
clf1 = CalibratedClassifierCV(clf1, cv=3)

# Classifier 2: SVC with best params
clf2 = SVC(C=10, kernel='rbf', random_state=42, class_weight='balanced',
    ↪probability=True)
clf2.fit(tr_X,tr_y)
```

```

clf2 = CalibratedClassifierCV(clf2, cv=3)

# Classifier 3: XGBoost with best params
clf3 = XGBClassifier(max_depth=2, n_estimators=300, scale_pos_weight=0.5)
clf3.fit(tr_X, tr_y)
clf3 = CalibratedClassifierCV(clf3, cv=3)

# Classifier 4: RF with best params
clf4 = RandomForestClassifier(max_depth=2, n_estimators=300)
clf4.fit(tr_X, tr_y)
clf4 = CalibratedClassifierCV(clf4, cv=3)

# Stack Classifier
sclf = StackingClassifier(classifiers=[clf1, clf2, clf3, clf4],
    ↪ meta_classifier=clf1, use_probas=True)

# Fit the model
sclf.fit(tr_X, tr_y)

```

```

[93]: StackingClassifier(average_probas=False,
    classifiers=[CalibratedClassifierCV(base_estimator=LogisticRegression(C=0.1,
        class_weight='balanced',
        dual=False,
        fit_intercept=True,
        intercept_scaling=1,
        l1_ratio=None,
        max_iter=100,
        multi_class='auto',
        n_jobs=None,
        penalty='l1',
        random_state=42,
        solver='liblinear',
        tol=0.0001,
        verbose=0,
        warm_start=False),
        cv=3, method='si...
    meta_classifier=CalibratedClassifierCV(base_estimator=LogisticRegression(C=0.1,
        class_weight='balanced',
        dual=False,
        fit_intercept=True,
        intercept_scaling=1,
        l1_ratio=None,
        max_iter=100,
        multi_class='auto',
        n_jobs=None,
        penalty='l1',
        random_state=42,

```

```

solver='liblinear',
tol=0.0001,
verbose=0,
warm_start=False),

cv=3,
method='sigmoid'),
store_train_meta_features=False, use_clones=True,
use_features_in_secondary=False, use_probas=True, verbose=0)

```

## 22 5.6.1 Kaggle score without top features

```

[94]: # Create a submission file format to submit in Kaggle
temp_id = df_test['id']
sclf_csv = sclf.predict_proba(ts_X)[: ,1]
sclf_df = pd.DataFrame(np.column_stack((temp_id,sclf_csv)),
    ↪ columns=['id', 'target'])
sclf_df['id'] = sclf_df['id'].astype('int32')
sclf_df.to_csv(data_dir+'/submission_sclf.csv', index=False)

```

```


[95]: image = plt.imread(data_dir+'/submission_sclf.png')
plt.figure(figsize=(18,5))
plt.imshow(image)

```

```

[95]: <matplotlib.image.AxesImage at 0x266f2521408>

```



## 23 5.6.2 Using Top features

```

[97]: # Fit the model
sclf.fit(tr_X[:,comb_top_feat], tr_y)

```

```

[97]: StackingClassifier(average_probas=False,
classifiers=[CalibratedClassifierCV(base_estimator=LogisticRegression(C=0.1,
    class_weight='balanced',
    dual=False,
    fit_intercept=True,
    intercept_scaling=1,
    l1_ratio=None,

```



```

        max_iter=100,
        multi_class='auto',
        n_jobs=None,
        penalty='l1',
        random_state=42,
        solver='liblinear',
        tol=0.0001,
        verbose=0,
        warm_start=False),

                                cv=3, method='si...
meta_classifier=CalibratedClassifierCV(base_estimator=LogisticRegression(C=0.1,
    class_weight='balanced',
    dual=False,
    fit_intercept=True,
    intercept_scaling=1,
    l1_ratio=None,
    max_iter=100,
    multi_class='auto',
    n_jobs=None,
    penalty='l1',
    random_state=42,
    solver='liblinear',
    tol=0.0001,
    verbose=0,
    warm_start=False),

                                cv=3,
                                method='sigmoid'),
    store_train_meta_features=False, use_clones=True,
    use_features_in_secondary=False, use_proba=True, verbose=0)

```

```

[98]: # Create a submission file format to submit in Kaggle
temp_id = df_test['id']
sclf_csv = sclf.predict_proba(ts_X[:, comb_top_feat])[:,1]
sclf_df = pd.DataFrame(np.column_stack((temp_id,sclf_csv)),  
    ↪columns=['id', 'target'])
sclf_df['id'] = sclf_df['id'].astype('int32')
sclf_df.to_csv(data_dir+'/submission_sclf_topfeat.csv', index=False)

```

```


[99]: image = plt.imread(data_dir+'/sclf_topfeat.png')
plt.figure(figsize=(18,5))
plt.imshow(image)

```

```

[99]: <matplotlib.image.AxesImage at 0x266f23771c8>

```



## 24 5.7 Voting Classifier (Without Stack Classifier + no weights)

```
[100]: # Import Voting Classifier
from mlxtend.classifier import EnsembleVoteClassifier
```

```
[102]: # Voting Classifier (See Docs: http://rasbt.github.io/mlxtend/user\_guide/classifier/EnsembleVoteClassifier/)
eclf = EnsembleVoteClassifier(clfs=[clf1, clf2, clf3, clf4])
# Fit the train data
eclf.fit(tr_X, tr_y)
```

```
[102]: EnsembleVoteClassifier(clfs=[CalibratedClassifierCV(base_estimator=LogisticRegression(C=0.1,
    class_weight='balanced',
    dual=False,
    fit_intercept=True,
    intercept_scaling=1,
    l1_ratio=None,
    max_iter=100,
    multi_class='auto',
    n_jobs=None,
    penalty='l1',
    random_state=42,
    solver='liblinear',
    tol=0.0001,
    verbose=0,
    warm_start=False),
    cv=3, method='sigmoid'),
    CalibratedClassi...
    max_depth=2,
    max_features='auto',
    max_leaf_nodes=None,
    max_samples=None,
    min_impurity_decrease=0.0,
    min_impurity_split=None,
    min_samples_leaf=1,
    min_samples_split=2,
    min_weight_fraction_leaf=0.0,
    n_estimators=300,
```

```

n_jobs=None,
oob_score=False,
random_state=None,
verbose=0,
warm_start=False),

cv=3, method='sigmoid']],
refit=True, verbose=0, voting='hard', weights=None)

```

## 25 5.7.1 Kaggle Score without top features

```

[105]: # Create a submission file format to submit in Kaggle
temp_id = df_test['id']
eclf_csv = eclf.predict_proba(ts_X)[:,:1]
eclf_df = pd.DataFrame(np.column_stack((temp_id,eclf_csv)),
    ↪columns=['id','target'])
eclf_df['id'] = eclf_df['id'].astype('int32')
eclf_df.to_csv(data_dir+'/submission_eclf.csv', index=False)


```

```

[106]: image = plt.imread(data_dir+'/submission_eclf.png')
plt.figure(figsize=(18,5))
plt.imshow(image)

```

[106]: <matplotlib.image.AxesImage at 0x266f23cd488>



## 26 5.7.2 Kaggle Score using top features

```

[107]: # Fit the model
eclf.fit(tr_X[:,comb_top_feat], tr_y)

```

```

[107]: EnsembleVoteClassifier(clfs=[CalibratedClassifierCV(base_estimator=LogisticRegression(C=0.1,
    class_weight='balanced',
    dual=False,
    fit_intercept=True,
    intercept_scaling=1,
    l1_ratio=None,
    max_iter=100,
    multi_class='auto',

```

```

n_jobs=None,
penalty='l1',
random_state=42,
solver='liblinear',
tol=0.0001,
verbose=0,
warm_start=False),

                                cv=3, method='sigmoid'),

                                CalibratedClassi...

                                max_depth=2,
                                max_features='auto',
                                max_leaf_nodes=None,
                                max_samples=None,
                                min_impurity_decrease=0.0,
                                min_impurity_split=None,
                                min_samples_leaf=1,
                                min_samples_split=2,
                                min_weight_fraction_leaf=0.0,
                                n_estimators=300,
                                n_jobs=None,
                                oob_score=False,
                                random_state=None,
                                verbose=0,
                                warm_start=False),

                                cv=3, method='sigmoid')],
                                refit=True, verbose=0, voting='hard', weights=None)

```

```

[108]: # Create a submission file format to submit in Kaggle
temp_id = df_test['id']
eclf_csv = eclf.predict_proba(ts_X[:,comb_top_feat])[:,1]
eclf_df = pd.DataFrame(np.column_stack((temp_id,eclf_csv)),
    ↪columns=['id', 'target'])
eclf_df['id'] = eclf_df['id'].astype('int32')
eclf_df.to_csv(data_dir+'/submission_eclf_topfeat.csv', index=False)

```

```


[109]: image = plt.imread(data_dir+'/eclf_topfeat.png')
plt.figure(figsize=(18,5))
plt.imshow(image)

```

```

[109]: <matplotlib.image.AxesImage at 0x266f16b9ac8>

```



## 27 5.8 Voting Classifier (With Stack Classifier + no weights)

```
[110]: # Voting Classifier (See Docs: http://rasbt.github.io/mlxtend/user\_guide/classifier/EnsembleVoteClassifier/)
eclf = EnsembleVoteClassifier(clfs=[clf1, clf2, clf3, clf4, sclf])
# Fit the train data
eclf.fit(tr_X, tr_y)
```

```
[110]: EnsembleVoteClassifier(clfs=[CalibratedClassifierCV(base_estimator=LogisticRegression(C=0.1,
    class_weight='balanced',
    dual=False,
    fit_intercept=True,
    intercept_scaling=1,
    l1_ratio=None,
    max_iter=100,
    multi_class='auto',
    n_jobs=None,
    penalty='l1',
    random_state=42,
    solver='liblinear',
    tol=0.0001,
    verbose=0,
    warm_start=False),
    cv=3, method='sigmoid'),
    CalibratedClassifierCV(
        fit_intercept=True,
        intercept_scaling=1,
        l1_ratio=None,
        max_iter=100,
        multi_class='auto',
        n_jobs=None,
        penalty='l1',
        random_state=42,
        solver='liblinear',
        tol=0.0001,
        verbose=0,
        warm_start=False),
    cv=3,
    method='sigmoid'),
    store_train_meta_features=False,
    use_clones=True,
    use_features_in_secondary=False,
    use_probabilities=True, verbose=0)]],
```


```
refit=True, verbose=0, voting='hard', weights=None)
```

## 28 5.8.1 Kaggle Score without top features

```
[111]: # Create a submission file format to submit in Kaggle
temp_id = df_test['id']
eclf_csv = eclf.predict_proba(ts_X)[: ,1]
eclf_df = pd.DataFrame(np.column_stack((temp_id,eclf_csv)),
    ↪columns=['id', 'target'])
eclf_df['id'] = eclf_df['id'].astype('int32')
eclf_df.to_csv(data_dir+'/submission_eclf_stack.csv', index=False)
```

```
[112]: image = plt.imread(data_dir+'/submission_eclf_stack.png')
plt.figure(figsize=(18,5))
plt.imshow(image)
```

```
[112]: <matplotlib.image.AxesImage at 0x266f1fdfb88>
```



## 29 5.8.2 Kaggle Score using top features

```
[114]: # Fit the model
eclf.fit(tr_X[:,comb_top_feat], tr_y)
```

```
[114]: EnsembleVoteClassifier(clfs=[CalibratedClassifierCV(base_estimator=LogisticRegression(C=0.1,
    class_weight='balanced',
    dual=False,
    fit_intercept=True,
    intercept_scaling=1,
    l1_ratio=None,
    max_iter=100,
    multi_class='auto',
    n_jobs=None,
    penalty='l1',
    random_state=42,
    solver='liblinear',
    tol=0.0001,
```

```

verbose=0,
warm_start=False),

                                cv=3, method='sigmoid'),
                                CalibratedClassi...
                                fit_intercept=True,
                                intercept_scaling=1,
                                l1_ratio=None,
                                max_iter=100,
                                multi_class='auto',
                                n_jobs=None,
                                penalty='l1',
                                random_state=42,
                                solver='liblinear',
                                tol=0.0001,
                                verbose=0,
                                warm_start=False),

                                cv=3,
                                method='sigmoid'),

                                store_train_meta_features=False,
                                use_clones=True,
                                use_features_in_secondary=False,
                                use_probabilities=True, verbose=0)],
                                refit=True, verbose=0, voting='hard', weights=None)

```

```

[115]: # Create a submission file format to submit in Kaggle
temp_id = df_test['id']
eclf_csv = eclf.predict_proba(ts_X[:,comb_top_feat])[:,1]
eclf_df = pd.DataFrame(np.column_stack((temp_id,eclf_csv)),
    columns=['id','target'])
eclf_df['id'] = eclf_df['id'].astype('int32')
eclf_df.to_csv(data_dir+'/submission_eclf_stack_topfeat.csv', index=False)

```

```


[116]: image = plt.imread(data_dir+'/submission_eclf_stack_topfeat.png')
plt.figure(figsize=(18,5))
plt.imshow(image)

```

```

[116]: <matplotlib.image.AxesImage at 0x266f1fb2608>

```



## 30 5.9 Voting Classifier (without Stack Classifier + weights)

```
[141]: # Voting Classifier (See Docs: http://rasbt.github.io/mlxtend/user\_guide/  
→ classifier/EnsembleVoteClassifier/)  
ecclf = EnsembleVoteClassifier(clfs=[clf1,clf2,clf3,clf4], weights=[0.4,0.2,0.  
→ 2,0.2])  
# Fit the train data  
ecclf.fit(tr_X,tr_y)
```

```
[141]: EnsembleVoteClassifier(clfs=[CalibratedClassifierCV(base_estimator=LogisticRegre  
ssion(C=0.1,  
    class_weight='balanced',  
    dual=False,  
    fit_intercept=True,  
    intercept_scaling=1,  
    l1_ratio=None,  
    max_iter=100,  
    multi_class='auto',  
    n_jobs=None,  
    penalty='l1',  
    random_state=42,  
    solver='liblinear',  
    tol=0.0001,  
    verbose=0,  
    warm_start=False),  
                                cv=3, method='sigmoid'),  
                                CalibratedClassi...  
    max_features='auto',  
    max_leaf_nodes=None,  
    max_samples=None,  
    min_impurity_decrease=0.0,  
    min_impurity_split=None,  
    min_samples_leaf=1,  
    min_samples_split=2,  
    min_weight_fraction_leaf=0.0,  
    n_estimators=300,  
    n_jobs=None,  
    oob_score=False,  
    random_state=None,  
    verbose=0,  
    warm_start=False),  
                                cv=3, method='sigmoid')],  
    refit=True, verbose=0, voting='hard',  
    weights=[0.4, 0.2, 0.2, 0.2])
```


# 5.9.1 Kaggle Score without top features



```
[120]: # Create a submission file format to submit in Kaggle
temp_id = df_test['id']
eclf_csv = eclf.predict_proba(ts_X)[: ,1]
eclf_df = pd.DataFrame(np.column_stack((temp_id,eclf_csv)),
    ↳ columns=['id', 'target'])
eclf_df['id'] = eclf_df['id'].astype('int32')
eclf_df.to_csv(data_dir+'/submission_eclf_weights.csv', index=False)
```

```
[121]: image = plt.imread(data_dir+'/submission_eclf_weights.png')
plt.figure(figsize=(18,5))
plt.imshow(image)
```

```
[121]: <matplotlib.image.AxesImage at 0x266f2f91d08>
```



## 31 5.9.2 Kaggle Score using top features

```
[142]: # Voting Classifier (See Docs: http://rasbt.github.io/mlxtend/user\_guide/classifier/EnsembleVoteClassifier/)
eclf = EnsembleVoteClassifier(clfs=[clf1,clf2,clf3,clf4], weights=[0.4,0.2,0.
    ↳ 2,0.2])
# Fit the train data
eclf.fit(tr_X[:,comb_top_feat],tr_y)
```

```
[142]: EnsembleVoteClassifier(clfs=[CalibratedClassifierCV(base_estimator=LogisticRegression(C=0.1,
    class_weight='balanced',
    dual=False,
    fit_intercept=True,
    intercept_scaling=1,
    l1_ratio=None,
    max_iter=100,
    multi_class='auto',
    n_jobs=None,
    penalty='l1',
    random_state=42,
    solver='liblinear',
    tol=0.0001,
    verbose=0,
```

```

warm_start=False),

                                cv=3, method='sigmoid'),

                                CalibratedClassi...
                                max_features='auto',
                                max_leaf_nodes=None,
                                max_samples=None,
                                min_impurity_decrease=0.0,
                                min_impurity_split=None,
                                min_samples_leaf=1,
                                min_samples_split=2,
                                min_weight_fraction_leaf=0.0,
                                n_estimators=300,
                                n_jobs=None,
                                oob_score=False,
                                random_state=None,
                                verbose=0,
                                warm_start=False),

                                cv=3, method='sigmoid')],
                                refit=True, verbose=0, voting='hard',
                                weights=[0.4, 0.2, 0.2, 0.2])

```

```

[143]: # Create a submission file format to submit in Kaggle
temp_id = df_test['id']
eclf_csv = eclf.predict_proba(ts_X[:,comb_top_feat])[:,1]
eclf_df = pd.DataFrame(np.column_stack((temp_id,eclf_csv)),
    ↪columns=['id', 'target'])
eclf_df['id'] = eclf_df['id'].astype('int32')
eclf_df.to_csv(data_dir+'/submission_eclf_weights_topfeat.csv', index=False)

```

```


[145]: image = plt.imread(data_dir+'/submission_eclf_weights_topfeat.png')
plt.figure(figsize=(18,5))
plt.imshow(image)

```

```

[145]: <matplotlib.image.AxesImage at 0x266f4cacf08>

```



## 32 5.10 Voting Classifier (with Stack Classifier + weights)

```
[135]: # Voting Classifier (See Docs: http://rasbt.github.io/mlxtend/user\_guide/classifier/EnsembleVoteClassifier/)
eclf = EnsembleVoteClassifier(clfs=[clf1,clf2,clf3,clf4,sclf], weights=[0.4,0.
    ↪1,0.1,0.2,0.2])
# Fit the train data
eclf.fit(tr_X,tr_y)
```


```
[135]: EnsembleVoteClassifier(clfs=[CalibratedClassifierCV(base_estimator=LogisticRegression(C=0.1,
    class_weight='balanced',
    dual=False,
    fit_intercept=True,
    intercept_scaling=1,
    l1_ratio=None,
    max_iter=100,
    multi_class='auto',
    n_jobs=None,
    penalty='l1',
    random_state=42,
    solver='liblinear',
    tol=0.0001,
    verbose=0,
    warm_start=False),
                                cv=3, method='sigmoid'),
                                CalibratedClassifierCV(
                                    intercept_scaling=1,
                                    l1_ratio=None,
                                    max_iter=100,
                                    multi_class='auto',
                                    n_jobs=None,
                                    penalty='l1',
                                    random_state=42,
                                    solver='liblinear',
                                    tol=0.0001,
                                    verbose=0,
                                    warm_start=False),
                                cv=3,
                                method='sigmoid'),
                                store_train_meta_features=False,
                                use_clones=True,
                                use_features_in_secondary=False,
                                use_probabilities=True, verbose=0)],
    refit=True, verbose=0, voting='hard',
    weights=[0.4, 0.1, 0.1, 0.2, 0.2])
```

### 33 5.10.1 Kaggle Score without top features

```
[136]: # Create a submission file format to submit in Kaggle
temp_id = df_test['id']
eclf_csv = eclf.predict_proba(ts_X)[: ,1]
eclf_df = pd.DataFrame(np.column_stack((temp_id,eclf_csv)),
    ↪columns=['id', 'target'])
eclf_df['id'] = eclf_df['id'].astype('int32')
eclf_df.to_csv(data_dir+'/submission_eclf_stack_weights.csv', index=False)
```

```
[137]: image = plt.imread(data_dir+'/submission_eclf_stack_weights.png')
plt.figure(figsize=(18,5))
plt.imshow(image)
```

```
[137]: <matplotlib.image.AxesImage at 0x266f272d048>
```



### 34 5.10.2 Kaggle Score using top features

```
[138]: # Voting Classifier (See Docs: http://rasbt.github.io/mlxtend/user\_guide/classifier/EnsembleVoteClassifier/)
eclf = EnsembleVoteClassifier(clfs=[clf1,clf2,clf3,clf4,sclf], weights=[0.4,0.
    ↪1,0.1,0.2,0.2])
# Fit the train data
eclf.fit(tr_X[:,comb_top_feat],tr_y)
```

```
[138]: EnsembleVoteClassifier(clfs=[CalibratedClassifierCV(base_estimator=LogisticRegression(C=0.1,
    class_weight='balanced',
    dual=False,
    fit_intercept=True,
    intercept_scaling=1,
    l1_ratio=None,
    max_iter=100,
    multi_class='auto',
    n_jobs=None,
    penalty='l1',
    random_state=42,
    solver='liblinear',
```

```

tol=0.0001,
verbose=0,
warm_start=False),

                                cv=3, method='sigmoid'),

                                CalibratedClassi...
                                intercept_scaling=1,
                                l1_ratio=None,
                                max_iter=100,
                                multi_class='auto',
                                n_jobs=None,
                                penalty='l1',
                                random_state=42,
                                solver='liblinear',
                                tol=0.0001,
                                verbose=0,
                                warm_start=False),

                                cv=3,
                                method='sigmoid'),

                                store_train_meta_features=False,
                                use_clones=True,
                                use_features_in_secondary=False,
                                use_probabilities=True, verbose=0)],

                                refit=True, verbose=0, voting='hard',
                                weights=[0.4, 0.1, 0.1, 0.2, 0.2])

```

```

[139]: # Create a submission file format to submit in Kaggle
temp_id = df_test['id']
eclf_csv = eclf.predict_proba(ts_X[:,comb_top_feat])[:,1]
eclf_df = pd.DataFrame(np.column_stack((temp_id,eclf_csv)),
    ↪columns=['id','target'])
eclf_df['id'] = eclf_df['id'].astype('int32')
eclf_df.to_csv(data_dir+'/submission_eclf_stack_weights_topfeat.csv',
    ↪index=False)


```

```

[140]: image = plt.imread(data_dir+'/submission_eclf_stack_weights_topfeat.png')
plt.figure(figsize=(18,5))
plt.imshow(image)

```

[140]: <matplotlib.image.AxesImage at 0x266f3060b48>



## 35 6. Summary of All Models

```
[151]: from prettytable import PrettyTable
x = PrettyTable()
x.field_names = ['Model', 'Features', 'Hyperparameter', 'Test Score']
x.add_row(['knn', 'AF', r"{'algorithm': 'kd_tree', 'n_neighbors': 35}", 0.693])
x.add_row(['Logistic Regression', 'AF', r"{'C': 0.1, 'penalty': 'l1', 'solver': 'liblinear'", 0.828])
x.add_row(['Logistic Regression', 'top10', r"{'C': 0.1, 'penalty': 'l1', 'solver': 'liblinear'", 0.821])
x.add_row(['Logistic Regression', 'top20', r"{'C': 0.1, 'penalty': 'l1', 'solver': 'liblinear'", 0.822])
x.add_row(['SVC', 'AF', r"{'C': 10, 'kernel': 'rbf'", 0.720])
x.add_row(['SVC', 'top10', r"{'C': 10, 'kernel': 'rbf'", 0.689])
x.add_row(['RandomForest', 'AF', r"{'max_depth': 2, 'n_estimators': 300}", 0.745])
x.add_row(['RandomForest', 'top10', r"{'max_depth': 2, 'n_estimators': 300}", 0.784])
x.add_row(['XGBoost', 'AF', r"{'max_depth': 2, 'n_estimators': 300}", 0.765])
x.add_row(['XGBoost', 'top10', r"{'max_depth': 2, 'n_estimators': 300}", 0.646])
x.add_row(['Stacking Classifier', 'AF', '-', 0.741])
x.add_row(['Stacking Classifier', 'topfeatures', '-', 0.808])
x.add_row(['Voting Classifier(no stacking + no weights)', 'AF', "-", 0.820])
x.add_row(['Voting Classifier(no stacking + no weights)', 'topfeatures', "-", 0.834])
x.add_row(['Voting Classifier(stacking + no weights)', 'AF', "-", 0.814])
x.add_row(['Voting Classifier(stacking + no weights)', 'topfeatures', "-", 0.831])
x.add_row(['Voting Classifier(no stacking + weights)', 'AF', "-", 0.830])
x.add_row(['Voting Classifier(no stacking + weights)', 'topfeatures', "-", 0.839])
x.add_row(['Voting Classifier(stacking + weights)', 'AF', "-", 0.827])
x.add_row(['Voting Classifier(stacking + weights)', 'topfeatures', "-", 0.838])
print(x)
```

Model	Features	Hyperparameter	Test Score
knn	AF	{'algorithm': 'kd_tree', 'n_neighbors': 35}	0.693
Logistic Regression	AF	{'C': 0.1, 'penalty': 'l1', 'solver': 'liblinear'}	0.828
Logistic Regression	top10	{'C': 0.1, 'penalty': 'l1', 'solver': 'liblinear'}	0.821
Logistic Regression	top20	{'C': 0.1, 'penalty': 'l1', 'solver': 'liblinear'}	0.822
SVC	AF	{'C':	

10, 'kernel': 'rbf'}		0.72			
SVC				top10	{'C':
10, 'kernel': 'rbf'}		0.689			
RandomForest				AF	
{'max_depth': 2, 'n_estimators': 300}				0.745	
RandomForest				top10	
{'max_depth': 2, 'n_estimators': 300}				0.784	
XGBoost				AF	
{'max_depth': 2, 'n_estimators': 300}				0.765	
XGBoost				top10	
{'max_depth': 2, 'n_estimators': 300}				0.646	
Stacking Classifier				AF	
-		0.741			
Stacking Classifier				topfeatures	
-		0.808			
Voting Classifier(No stacking + no weights)				AF	
-		0.82			
Voting Classifier(No stacking + no weights)				topfeatures	
-		0.834			
Voting Classifier(stacking + no weights)				AF	
-		0.814			
Voting Classifier(stacking + no weights)				topfeatures	
-		0.831			
Voting Classifier(no stacking + weights)				AF	
-		0.83			
Voting Classifier(no stacking + weights)				topfeatures	
-		0.839			
Voting Classifier(stacking + weights)				AF	
-		0.827			
Voting Classifier(stacking + weights)				topfeatures	
-		0.838			
+-----+-----+-----+					
-----+-----+					

Notation: 1. AF: All features 2. top10: Find top 10 features using forward feature selections of that model 3. top20: Find top 20 features using forward feature selections of that model 4. topfeatures: combining all the top features generated using forward feature selection of that model.

[ ]: